

dm25s1

Topic 06 : Data Modelling

Part 02 : Column Encoding

Preparation

Data Handling

Exploring Data 1

Dr Bernard Butler
Department of Computing and Mathematics, WIT.
(bernard.butler@setu.ie)

Exploring Data 2

Building Models

Prediction

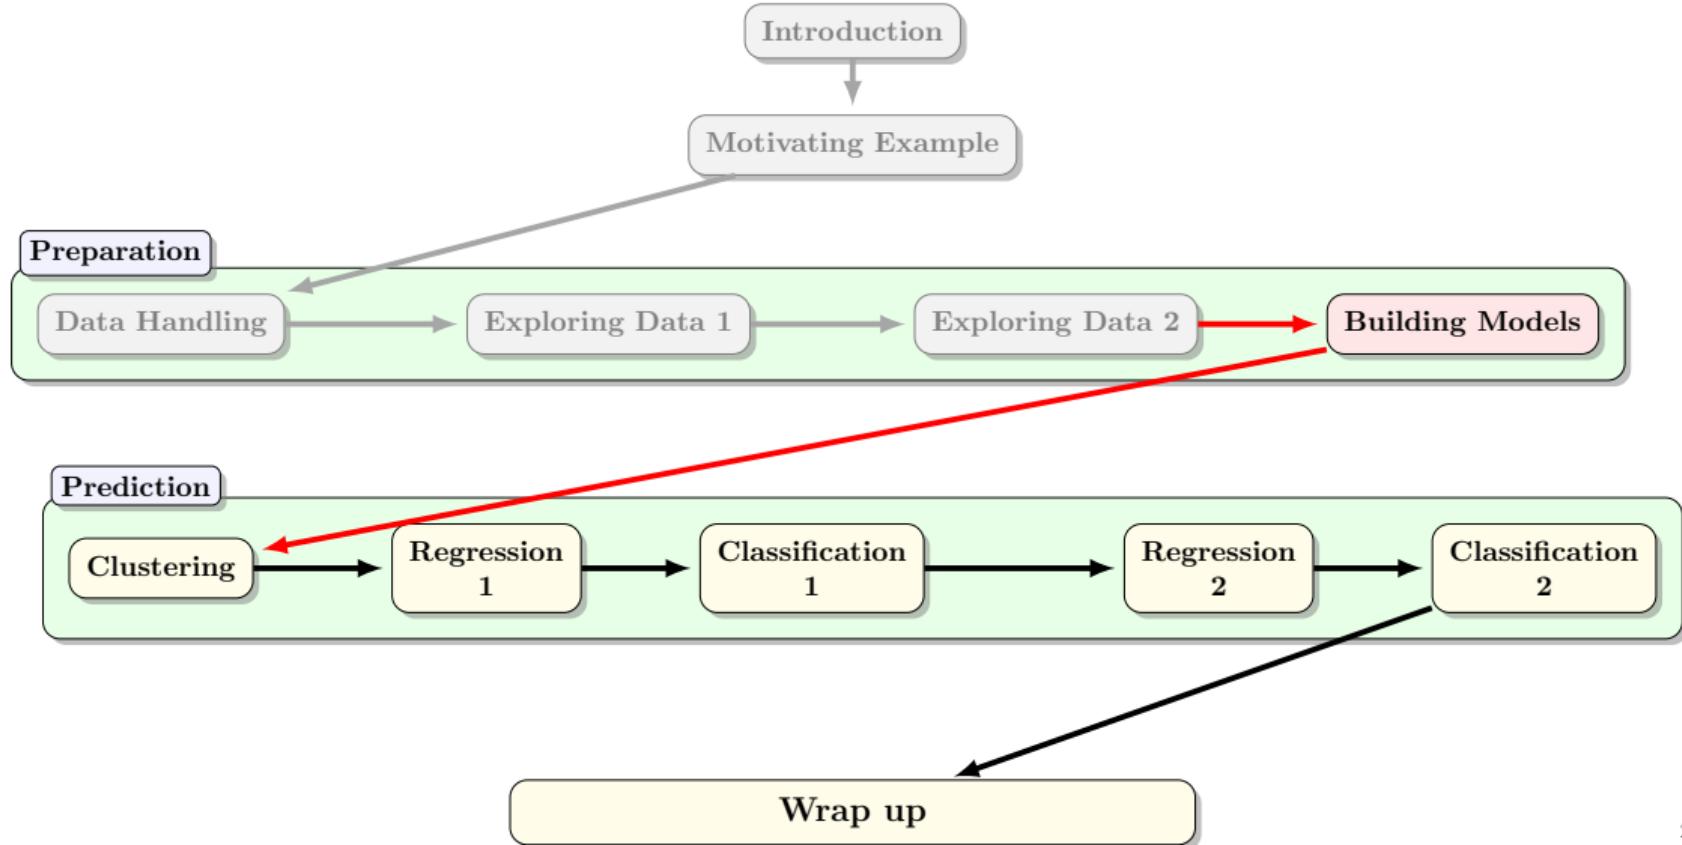
Autumn Semester, 2025

Outline

- Encoding categorical features, using pandas
- Looking ahead to feature preparation
- Overview of ML and what we have achieved

Wrap up

Data Mining (Week 6)



Outline

1. Modelling Process	3
1.1. Models and error	5
2. Categorical Columns	7
3. Wrap up	15
4. Resources	20

The Pipeline Metaphor

Model Building Pipeline



Defining the Goal



Building the Model



Interpreting the Model



Preparing the Dataset for ML



Evaluating the Model

Source: Dataiku

The Pipeline Metaphor

Model Building Pipeline



Defining the Goal



Building the Model



Interpreting the Model



Preparing the Dataset for ML




Evaluating the Model

Source: Dataiku


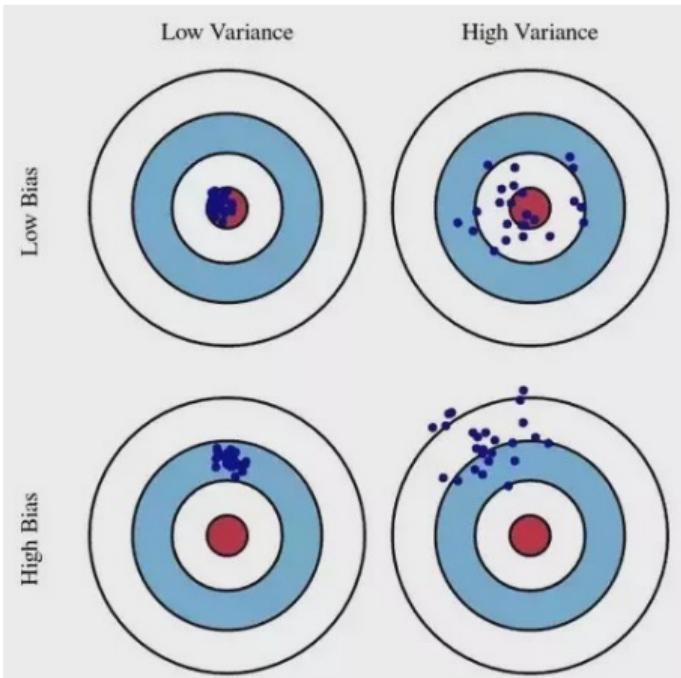
Comments

- We saw the first two stages in previous weeks
- This week we look at the remaining stages
- Of course this pipeline is a simplification. In reality it is iterative.


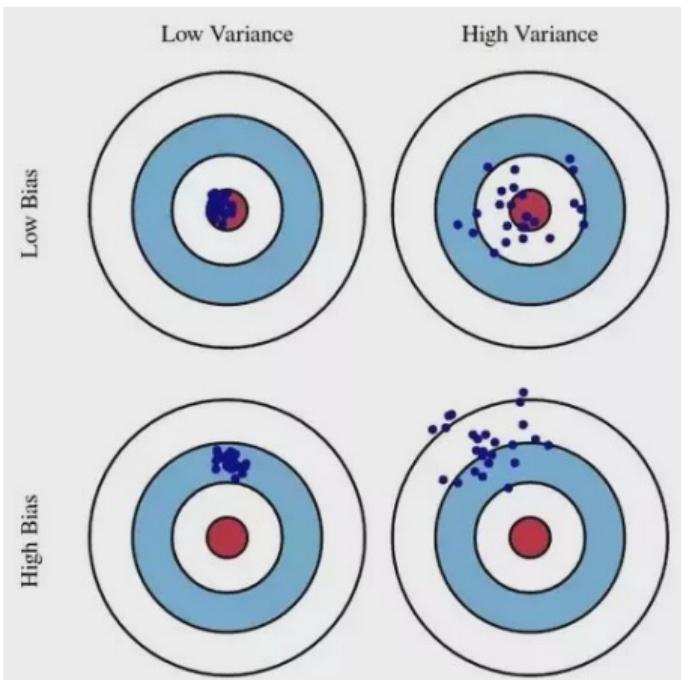
Bias-Variance and Total Error



Bias-Variance and Total Error



Bias-Variance and Total Error



Look for parameters α that minimise the generalization error (estimated using the test set that was not used during training)

Example Model Types

Model	Applications	Concerns
Logistic Regression	X-ray classification	Regression with transformed variable
Fully connected networks	Classification	Classical ANN: choose encoding and size
Convolutional Neural Networks	Image processing	deep learning - choose segmentation
Recurrent Neural Networks	Voice recognition	ANN with feedback - how much?
Random Forest	Fraud Detection	Ensemble method - how many?
Reinforcement Learning	Learning by trial and error	Choose goal and penalties
Generative Models	Text, Image creation	Choose parameters
K-means	Segmentation	Choose distance function and k
k-Nearest Neighbors	Recommendation systems	Choose distance function and k
Bayesian Classifiers	Spam and noise filtering	Deal with imbalances

Outline

1. Modelling Process	3
1.1. Models and error	5
2. Categorical Columns	7
3. Wrap up	15
4. Resources	20

Using Categorical Features in (Logistic) Regression

➤ How can Categorical-valued features participate in linear models?

Using Categorical Features in (Logistic) Regression

How can Categorical-valued features participate in linear models?

Given the following fragment of a dataset, where the goal is to predict the salary of employees in a large organisation:

```
df = pd.read_csv('data/team.csv',\n                 index_col="Name")  
df
```

	Role	Skilled	Salary
Name			
Alice	Designer	Yes	40000
Bob	Programmer	No	25000
Carol	Tester	No	30000

Using Categorical Features in (Logistic) Regression

How can Categorical-valued features participate in linear models?

Given the following fragment of a dataset, where the goal is to predict the salary of employees in a large organisation:

```
df = pd.read_csv('data/team.csv',  
                 index_col="Name")  
df
```

	Role	Skilled	Salary
Name			
Alice	Designer	Yes	40000
Bob	Programmer	No	25000
Carol	Tester	No	30000

How can this data be represented by a linear model, where all quantities must take numeric values?

Using pandas .getdummies() on a binary-valued column

```
dfSkilledDummies = pd.get_dummies(df['Skilled'], \
    prefix='Skilled', \
    dtype=int)
dfSkilledDummies
```

	Skilled_No	Skilled_Yes
Name		
Alice	0	1
Bob	1	0
Carol	1	0

Using pandas .getdummies() on a binary-valued column

```
dfSkilledDummies = pd.get_dummies(df['Skilled'],\n    prefix='Skilled',\n    dtype=int)\n\ndfSkilledDummies
```

Skilled_No **Skilled_Yes**

Name

Alice	0	1
Bob	1	0
Carol	1	0

➤ Note that a binary-valued column becomes 2 dummy columns

Reducing redundancy (by 1) in 2 dummy columns

```
dfSkilledIndicators = pd.get_dummies(df['Skilled'], \  
    prefix='Skilled', \  
    drop_first=True, \  
    dtype=int)\ \  
    .rename(columns={"Skilled_Yes": "IsSkilled"})  
dfSkilledIndicators
```

IsSkilled	
Name	
Alice	1
Bob	0
Carol	0

Reducing redundancy (by 1) in 2 dummy columns

```
dfSkilledIndicators = pd.get_dummies(df['Skilled'], \
    prefix='Skilled', \
    drop_first=True, \
    dtype=int)\ \
    .rename(columns={"Skilled_Yes": "IsSkilled"})  
dfSkilledIndicators
```

IsSkilled	
Name	
Alice	1
Bob	0
Carol	0

➤ A single indicator column can replace a group of 2 dummy columns

Using pandas .getdummies() on a multi-valued column

```
dfRoleDummies = pd.get_dummies(df['Role'], \
    prefix='Role', \
    dtype=int)
dfRoleDummies
```

	Role_Designer	Role_Programmer	Role_Tester
Name			
Alice	1	0	0
Bob	0	1	0
Carol	0	0	1

Using pandas .getdummies() on a multi-valued column

```
dfRoleDummies = pd.get_dummies(df['Role'], \
    prefix='Role', \
    dtype=int)
dfRoleDummies
```

	Role_Designer	Role_Programmer	Role_Tester
Name			
Alice	1	0	0
Bob	0	1	0
Carol	0	0	1

➤ Note that an n -valued column becomes n dummy columns

Reducing redundancy (by 1) in n dummy columns

```
dfRoleIndicators = pd.get_dummies(df['Role'], \
    prefix='Role', \
    drop_first=True, \
    dtype=int)\ \
    .rename(columns={\
        "Role_Programmer": "IsProgrammer", \
        "Role_Tester": "IsTester"})  
dfRoleIndicators
```

	IsProgrammer	IsTester
Name		
Alice	0	0
Bob	1	0
Carol	0	1

Reducing redundancy (by 1) in n dummy columns

```
dfRoleIndicators = pd.get_dummies(df['Role'], \
    prefix='Role', \
    drop_first=True, \
    dtype=int)\ \
    .rename(columns={\
        "Role_Programmer": "IsProgrammer", \
        "Role_Tester": "IsTester"})  
dfRoleIndicators
```

	IsProgrammer	IsTester
Name		
Alice	0	0
Bob	1	0
Carol	0	1

$\triangleright n - 1$ indicator columns can replace a group of n dummy columns

Deriving and using dummy/indicator features

- Identify potential categorical features in EDA Pass 1

Deriving and using dummy/indicator features

- Identify potential categorical features in EDA Pass 1
- Identify whether each feature is (potentially) *usable* in EDA Pass 2

Deriving and using dummy/indicator features

- Identify potential categorical features in EDA Pass 1
- Identify whether each feature is (potentially) *usable* in EDA Pass 2
- Identify whether each feature is (potentially) *useful* in EDA Pass 3

Deriving and using dummy/indicator features

- Identify potential categorical features in EDA Pass 1
- Identify whether each feature is (potentially) *usable* in EDA Pass 2
- Identify whether each feature is (potentially) *useful* in EDA Pass 3
- Add all potentially usable and useful features (regardless of type) to a list F

Deriving and using dummy/indicator features

- Identify potential categorical features in EDA Pass 1
- Identify whether each feature is (potentially) *usable* in EDA Pass 2
- Identify whether each feature is (potentially) *useful* in EDA Pass 3
- Add all potentially usable and useful features (regardless of type) to a list \mathcal{F}
- For each categorical feature f_j in \mathcal{F} having n levels

Deriving and using dummy/indicator features

- Identify potential categorical features in EDA Pass 1
- Identify whether each feature is (potentially) *usable* in EDA Pass 2
- Identify whether each feature is (potentially) *useful* in EDA Pass 3
- Add all potentially usable and useful features (regardless of type) to a list \mathcal{F}
- For each categorical feature f_j in \mathcal{F} having n levels
 - Derive $n - 1$ indicator features \tilde{f}_j^k , where $k = 1, \dots, n - 1$

Deriving and using dummy/indicator features

- Identify potential categorical features in EDA Pass 1
- Identify whether each feature is (potentially) *usable* in EDA Pass 2
- Identify whether each feature is (potentially) *useful* in EDA Pass 3
- Add all potentially usable and useful features (regardless of type) to a list \mathbf{F}
- For each categorical feature f_j in \mathbf{F} having n levels
 - Derive $n - 1$ indicator features \tilde{f}_j^k , where $k = 1, \dots, n - 1$
 - Replace the original categorical feature f_j in \mathbf{F} with the derived indicator features \tilde{f}_j^k .

Deriving and using dummy/indicator features

- Identify potential categorical features in EDA Pass 1
- Identify whether each feature is (potentially) *usable* in EDA Pass 2
- Identify whether each feature is (potentially) *useful* in EDA Pass 3
- Add all potentially usable and useful features (regardless of type) to a list \mathbf{F}
- For each categorical feature f_j in \mathbf{F} having n levels
 - Derive $n - 1$ indicator features \tilde{f}_j^k , where $k = 1, \dots, n - 1$
 - Replace the original categorical feature f_j in \mathbf{F} with the derived indicator features \tilde{f}_j^k .
- Build the model using the features in \mathbf{F} .

Analysis of pandas dummies

➤ Pandas enables us to convert unordered categorical features into sets of (0,1)-valued numeric features

Analysis of pandas dummies

➤ Pandas enables us to convert unordered categorical features into sets of (0,1)-valued numeric features

But what about...

Analysis of pandas dummies

➤ Pandas enables us to convert unordered categorical features into sets of (0,1)-valued numeric features

But what about...

- ➊ Ordered categorical features - can we do better than treating them as unordered (and losing information)?

Analysis of pandas dummies

➤ Pandas enables us to convert unordered categorical features into sets of (0,1)-valued numeric features

But what about...

- ➊ Ordered categorical features - can we do better than treating them as unordered (and losing information)?
- ➋ Categorical targets (whether ordered or not) - how should we handle these?

Analysis of pandas dummies

➤ Pandas enables us to convert unordered categorical features into sets of (0,1)-valued numeric features

But what about...

- ➊ Ordered categorical features - can we do better than treating them as unordered (and losing information)?
- ➋ Categorical targets (whether ordered or not) - how should we handle these?
- ➌ How do we handle data (that includes categorical columns) that is split into training and test?

Analysis of pandas dummies

► Pandas enables us to convert unordered categorical features into sets of (0,1)-valued numeric features


But what about...

- ➊ Ordered categorical features - can we do better than treating them as unordered (and losing information)?
- ➋ Categorical targets (whether ordered or not) - how should we handle these?
- ➌ How do we handle data (that includes categorical columns) that is split into training and test?
- ➍ How can we reverse the operation (i.e., return from (0,1)-valued columns to categorical columns)?

Outline


1. Modelling Process	3
1.1. Models and error	5
2. Categorical Columns	7
3. Wrap up	15
4. Resources	20

An overview of Machine Learning




On first glance, this might seem overwhelming, but note that the boxes are colour-coded, so that related

Where feature preparation fits...



Where feature preparation fits...



Feature Engineering more generally

- Use of pandas `getdummies()` requires care and is not always suitable

Feature Engineering more generally

- Use of pandas `getdummies()` requires care and is not always suitable
- Scikit-learn, imported as the `sklearn` package, supports a variety of column transformations

Feature Engineering more generally

- Use of pandas `getdummies()` requires care and is not always suitable
- Scikit-learn, imported as the `sklearn` package, supports a variety of column transformations
- Categorical columns can be features or targets, ordered or unordered

Feature Engineering more generally

- Use of pandas `getdummies()` requires care and is not always suitable
- Scikit-learn, imported as the `sklearn` package, supports a variety of column transformations
- Categorical columns can be features or targets, ordered or unordered
- Can also rescale numerical columns, or encode more exotic columns (other datatypes, computed columns, ...)

Feature Engineering more generally

- Use of pandas `getdummies()` requires care and is not always suitable
- Scikit-learn, imported as the `sklearn` package, supports a variety of column transformations
- Categorical columns can be features or targets, ordered or unordered
- Can also rescale numerical columns, or encode more exotic columns (other datatypes, computed columns, ...)
- As seen in the schematic, if an ML procedure is unsuccessful, more feature engineering should be considered - it can help.

Summary

- We have reviewed different types of models and considered their general form.

Summary

- We have reviewed different types of models and considered their general form.
- We looked at the goals of modelling: minimise predictive error.

Summary

- We have reviewed different types of models and considered their general form.
- We looked at the goals of modelling: minimise predictive error.
- We considered how feature engineering can help.

Summary

- We have reviewed different types of models and considered their general form.
- We looked at the goals of modelling: minimise predictive error.
- We considered how feature engineering can help.
 - Scaling numerical features, so that variation is treated fairly between features.

Summary

- We have reviewed different types of models and considered their general form.
- We looked at the goals of modelling: minimise predictive error.
- We considered how feature engineering can help.
 - Scaling numerical features, so that variation is treated fairly between features.
 - Choosing a subset of features (more to come in future weeks...), looking for the sweet spot between under- and over-fitting.

Summary

- We have reviewed different types of models and considered their general form.
- We looked at the goals of modelling: minimise predictive error.
- We considered how feature engineering can help.
 - Scaling numerical features, so that variation is treated fairly between features.
 - Choosing a subset of features (more to come in future weeks...), looking for the sweet spot between under- and over-fitting.
 - Encoding categorical features as numerical dummy features (more to come in future weeks...), so they can participate in linear models

Summary

- We have reviewed different types of models and considered their general form.
- We looked at the goals of modelling: minimise predictive error.
- We considered how feature engineering can help.
 - Scaling numerical features, so that variation is treated fairly between features.
 - Choosing a subset of features (more to come in future weeks...), looking for the sweet spot between under- and over-fitting.
 - Encoding categorical features as numerical dummy features (more to come in future weeks...), so they can participate in linear models
- In subsequent weeks we will put this theory into practice.

Outline

1. Modelling Process	3
1.1. Models and error	5
2. Categorical Columns	7
3. Wrap up	15
4. Resources	20

Resources

- **A Summary of the Basic Machine Learning Models**

towardsdatascience.com/a-summary-of-the-basic-machine-learning-models-e0a65627ecbe

- **Train-Test Split for Evaluating Machine Learning Algorithms**

<https://machinelearningmastery.com/>

[train-test-split-for-evaluating-machine-learning-algorithms](https://machinelearningmastery.com/train-test-split-for-evaluating-machine-learning-algorithms)

This week I have focused on the theory rather than its (python) implementation. This is a nice article that covers the implementation side of things.

- **Cross-Validation: Estimator Evaluator**

medium.com/swlh/cross-validation-estimator-evaluator-897d28afb4ff

Nice article that covers cross-validation in a lot more detail — we will be using many of these variants in later weeks, especially k-fold stratified.