

Data Mining (Week 1)

dm25s1

Topic 06 : Data Modelling

Part 02 : Column Encoding Scaling

Preparation

Data Handling

Motivating Example
Exploring Data 1

Dr Bernard Butler

Department of Computing and Mathematics, WIT
(bernard.butler@setu.ie)

Exploring Data 2

Building Models

Autumn Semester, 2025

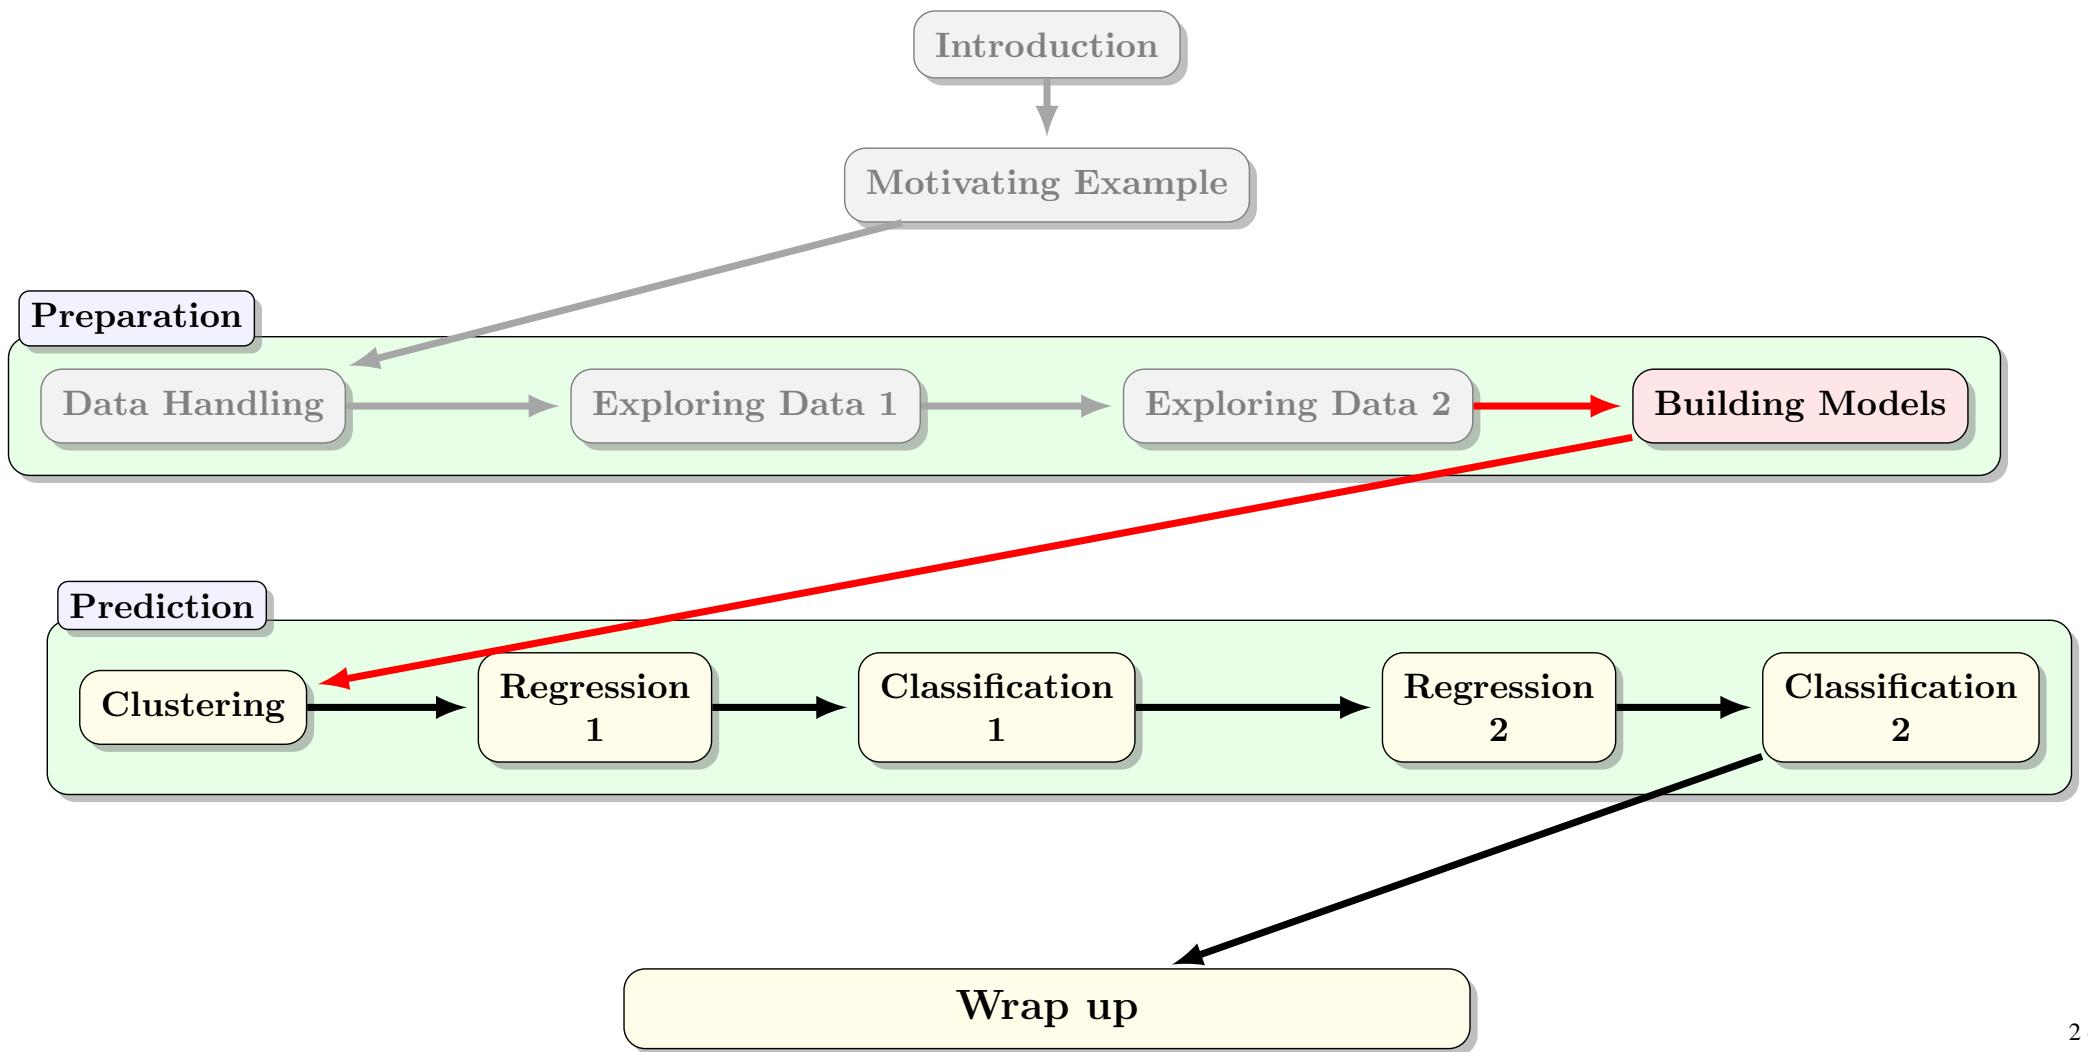
Prediction

Outline

- Encoding categorical features, using pandas
- Scaling numeric features
- Looking ahead to feature preparation
- Overview of ML and what we have achieved

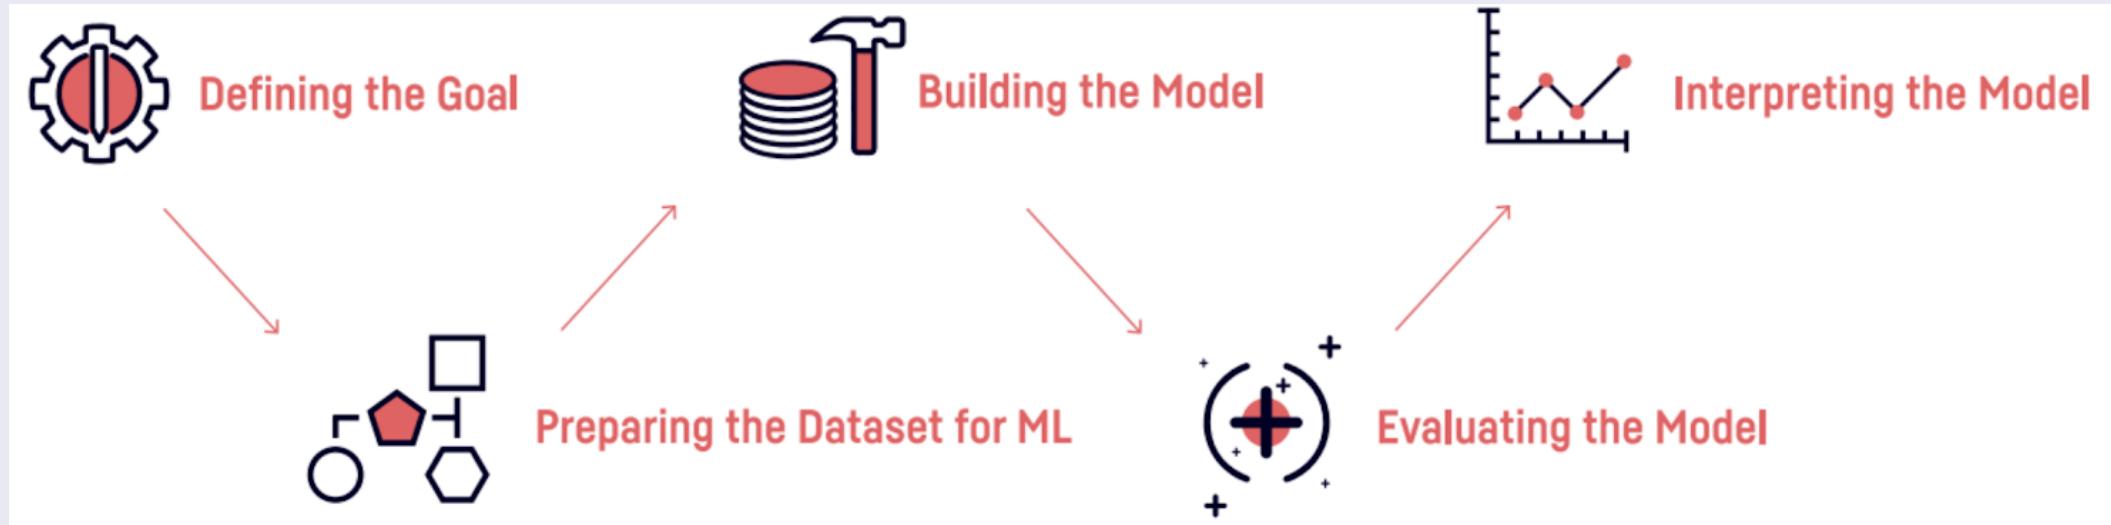
Wrap up

Data Mining (Week 6)



The Pipeline Metaphor

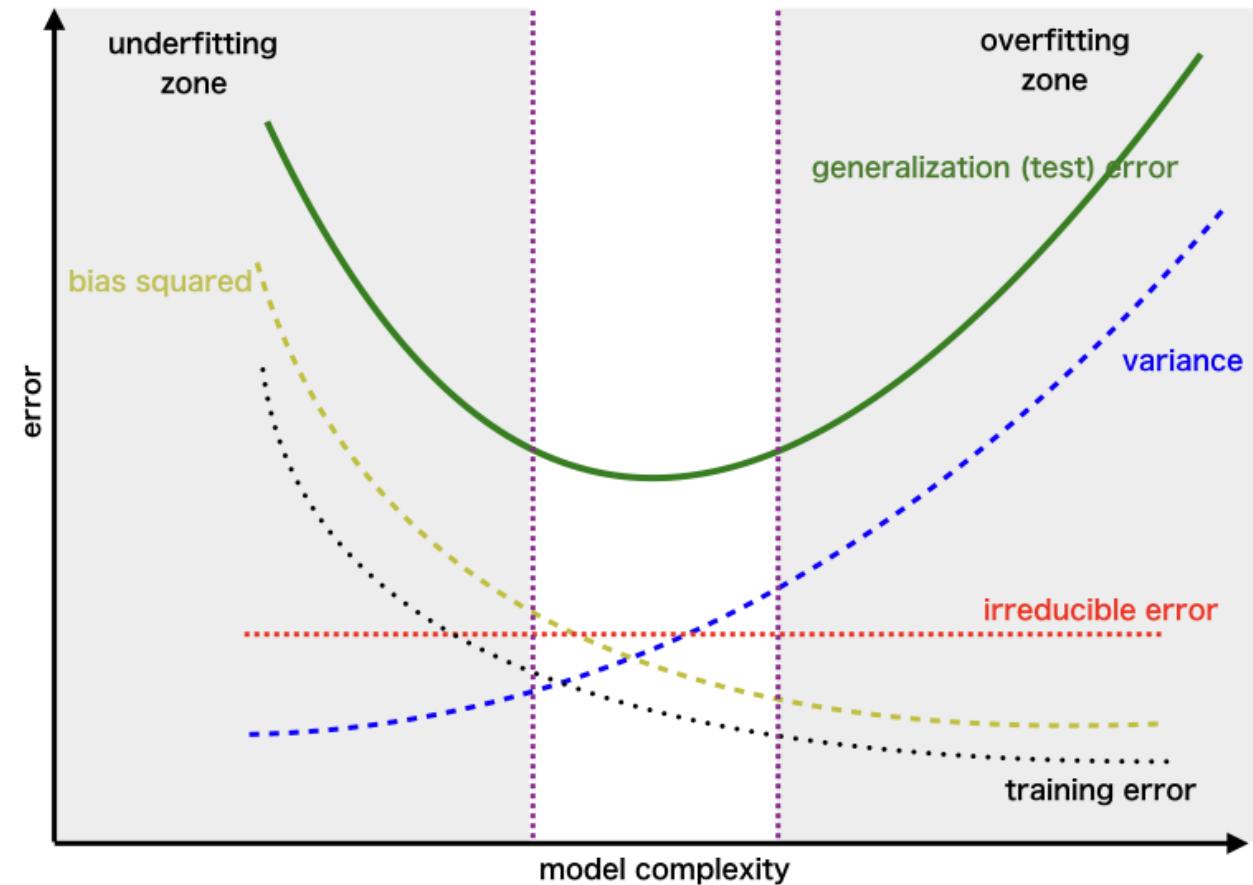
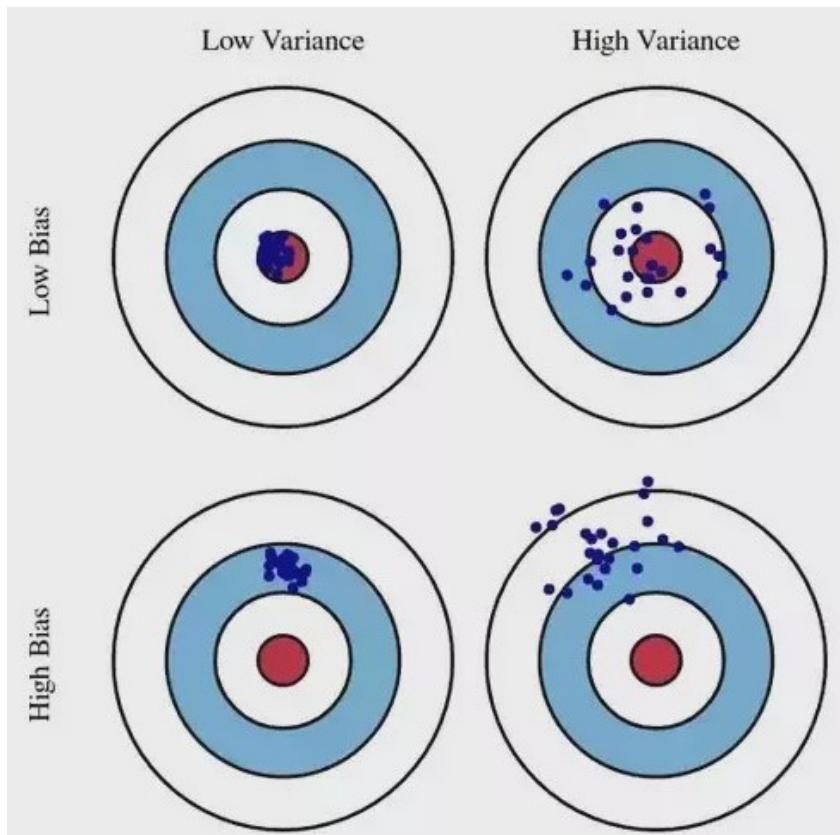
Model Building Pipeline



Comments

- We saw the first two stages in previous weeks
- This week we look at the remaining stages
- Of course this pipeline is a simplification. In reality it is iterative.

Bias-Variance and Total Error



Look for parameters a that minimise the generalization error (estimated using the test set that was not used during training)

Example Model Types

Model	Applications	Concerns
Logistic Regression	X-ray classification	Regression with transformed variable
Fully connected networks	Classification	Classical ANN: choose encoding and size
Convolutional Neural Networks	Image processing	deep learning - choose segmentation
Recurrent Neural Networks	Voice recognition	ANN with feedback - how much?
Random Forest	Fraud Detection	Ensemble method - how many?
Reinforcement Learning	Learning by trial and error	Choose goal and penalties
Generative Models	Text, Image creation	Choose parameters
K-means	Segmentation	Choose distance function and k
k-Nearest Neighbors	Recommendation systems	Choose distance function and k
Bayesian Classifiers	Spam and noise filtering	Deal with imbalances

Using Categorical Features in (Logistic) Regression

➤ How can Categorical-valued features participate in linear models?

Given the following fragment of a dataset, where the goal is to predict the salary of employees in a large organisation:

```
df = pd.read_csv('data/team.csv', \
                  index_col="Name")  
df
```

	Role	Skilled	Salary
	Name		
Alice	Designer	Yes	40000
Bob	Programmer	No	25000
Carol	Tester	No	30000

How can this data be represented by a linear model, where all quantities must take numeric values?

Using pandas .getdummies() on a binary-valued column

```
dfSkilledDummies = pd.get_dummies(df['Skilled'],  
prefix='Skilled',  
dtype=int)  
dfSkilledDummies
```

Skilled_No Skilled_Yes

Name

Alice 0

1

Bob 1

0

Carol 1

0

➤ Note that a binary-valued column becomes 2 dummy columns

Reducing redundancy (by 1) in 2 dummy columns

```
dfSkilledIndicators = pd.get_dummies(df['Skilled'], \
prefix='Skilled', \
drop_first=True, \
dtype=int)\ \
.rename(columns={"Skilled_Yes": "IsSkilled"})  
dfSkilledIndicators
```

IsSkilled	
Name	
Alice	1
Bob	0
Carol	0

➤ A single indicator column can replace a group of 2 dummy columns

Using pandas .getdummies() on a multi-valued column

```
dfRoleDummies = pd.get_dummies(df['Role'], \
    prefix='Role', \
    dtype=int)
dfRoleDummies
```

	Role_Designer	Role_Programmer	Role_Tester
Name			
Alice	1	0	0
Bob	0	1	0
Carol	0	0	1

➤ Note that an n -valued column becomes n dummy columns

Reducing redundancy (by 1) in n dummy columns

```
dfRoleIndicators = pd.get_dummies(df['Role'], \
prefix='Role', \
drop_first=True, \
dtype=int) \
.rename(columns={ \
    "Role_Programmer": "IsProgrammer", \
    "Role_Tester": "IsTester"})
```

dfRoleIndicators

	IsProgrammer	IsTester
Name		
Alice	0	0
Bob	1	0
Carol	0	1

➤ $n - 1$ indicator columns can replace a group of n dummy columns

Deriving and using dummy/indicator features

- Identify potential categorical features in EDA Pass 1
- Identify whether each feature is (potentially) *usable* in EDA Pass 2
- Identify whether each feature is (potentially) *useful* in EDA Pass 3
- Add all potentially usable and useful features (regardless of type) to a list F
- For each categorical feature f_j in F having n levels
 - Derive $n - 1$ indicator features \tilde{f}_j^k , where $k = 1, \dots, n - 1$
 - Replace the original categorical feature f_j in F with the derived indicator features \tilde{f}_j^k .
- Build the model using the features in F .

Analysis of pandas dummies

► Pandas enables us to convert unordered categorical features into sets of (0,1)-valued numeric features

But what about . . .

- ① Ordered categorical features - can we do better than treating them as unordered (and losing information)?
- ② Categorical targets (whether ordered or not) - how should we handle these?
- ③ How do we handle data (that includes categorical columns) that is split into training and test?
- ④ How can we reverse the operation (i.e., return from (0,1)-valued columns to categorical columns)?

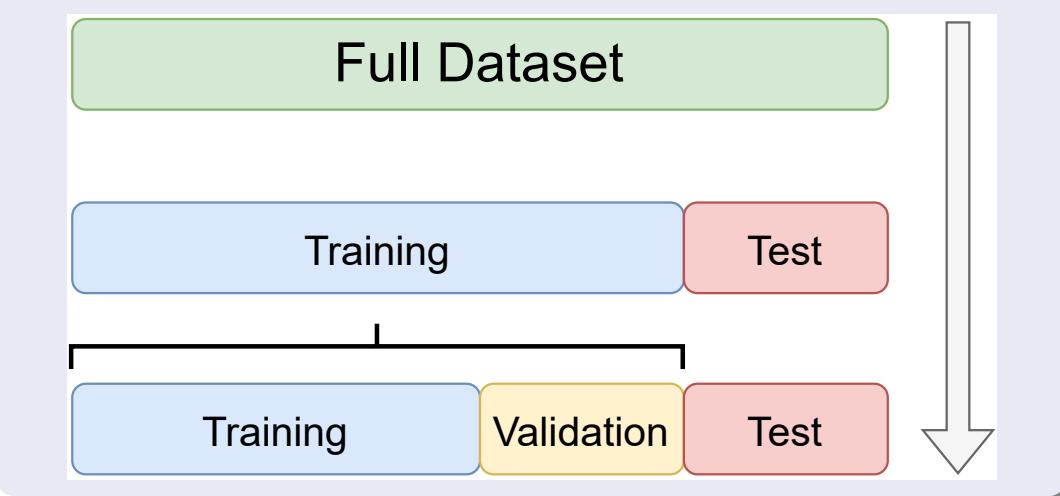
Training, test and valuation subsets: 3-way Holdout

Why Split?

Hold back some data to check how the model is doing.

- **Training** data is sample used to fit the model parameters.
- **Test** data is sample used to test the final model fitted to the training data.
- **Validation** data is sample used to test each interim model while tuning it.

Typical Splits

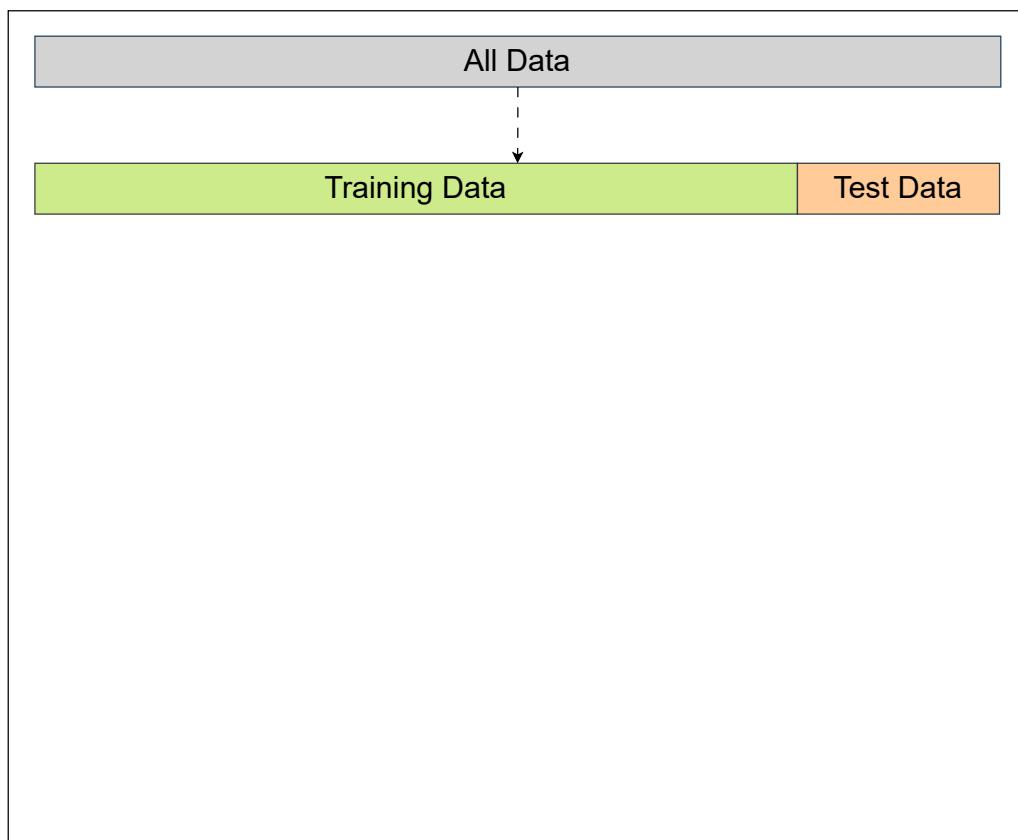


sklearn example

```

from sklearn.model_selection import train_test_split
trainVal, test = train_test_split(df, test_size=0.2, seed=42)
train, validation = train_test_split(trainVal, test_size=0.1)
  
```

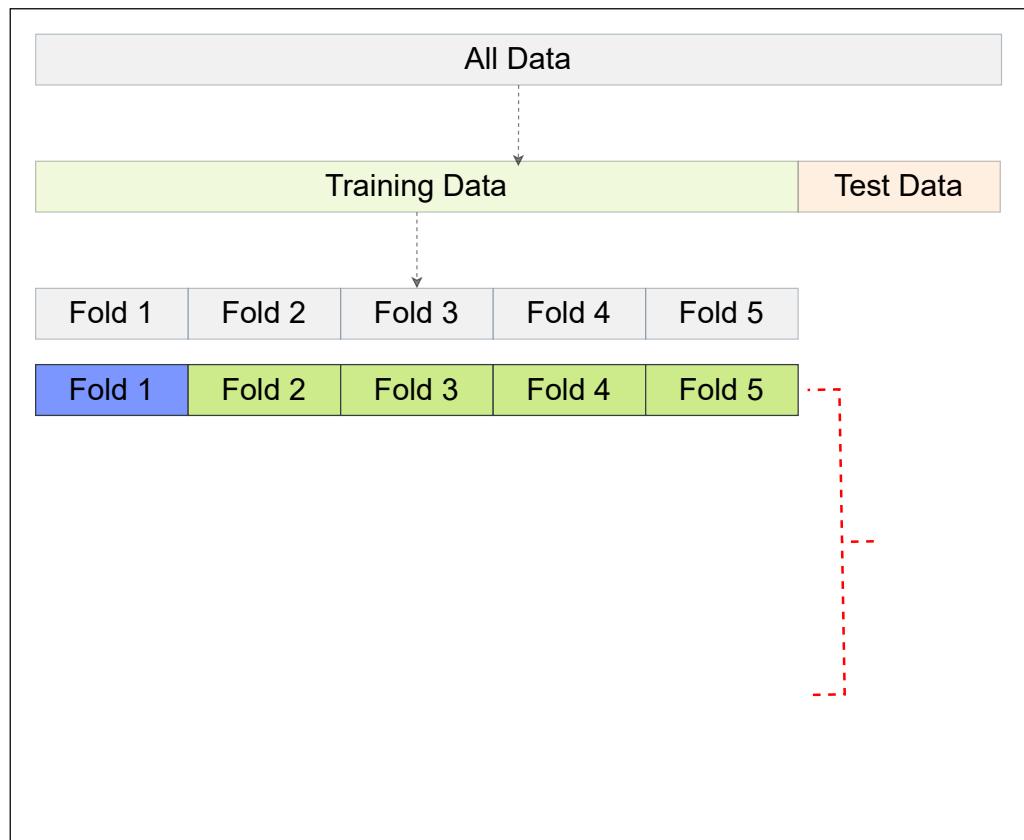
K-fold cross validation - preparation



Before Cross Validation

- Perform the train-test split
- Using the **Training Set** only
 - Perform 3-pass EDA
 - Encode the categorical features
 - Scale the numerical features
 - Other feature engineering steps...

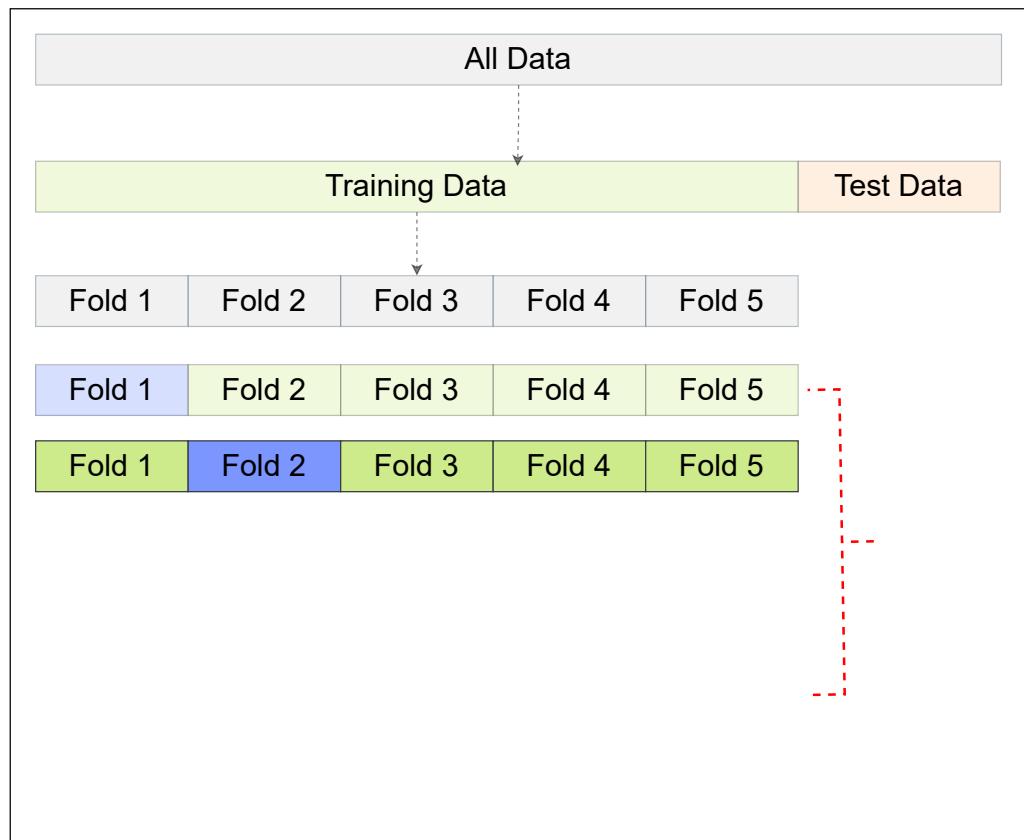
K-fold cross validation - initialise, pass 1



Initialise and perform first pass

- Randomly split the training set into K folds
- In the diagram, $K = 5$ but $K = 10$ is common
- In the first pass, set $i = 1$
- Fold $i = 1$ is held back for validation
- Folds $i = 2, \dots, K$ are used to train a model, giving parameters $\boldsymbol{a}^{(1)}$.
- The model is evaluated against fold $i = 1$
- The distance between the actual and predicted targets is calculated and recorded as $d^{(1)}$.

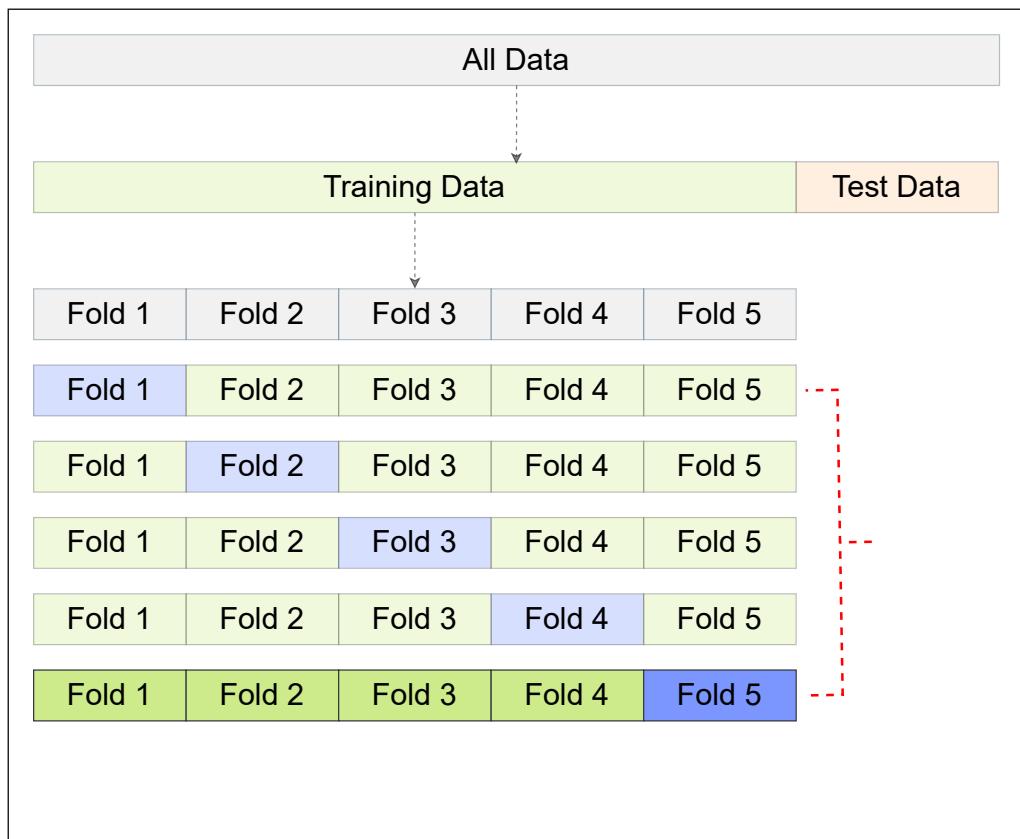
K-fold cross validation - initialise, pass 2



Perform second pass

- In the second pass, set $i = 2$
- Fold $i = 2$ is held back for validation
- Folds $i = 1, 3, \dots, K$ are used to train a model, giving parameters $\mathbf{a}^{(2)}$.
- The model is evaluated against fold $i = 2$
- The distance between the actual and predicted targets is calculated and recorded as $d^{(2)}$.
- The parameters $\mathbf{a}^{(2)}$ and distance $d^{(2)}$ are recorded.

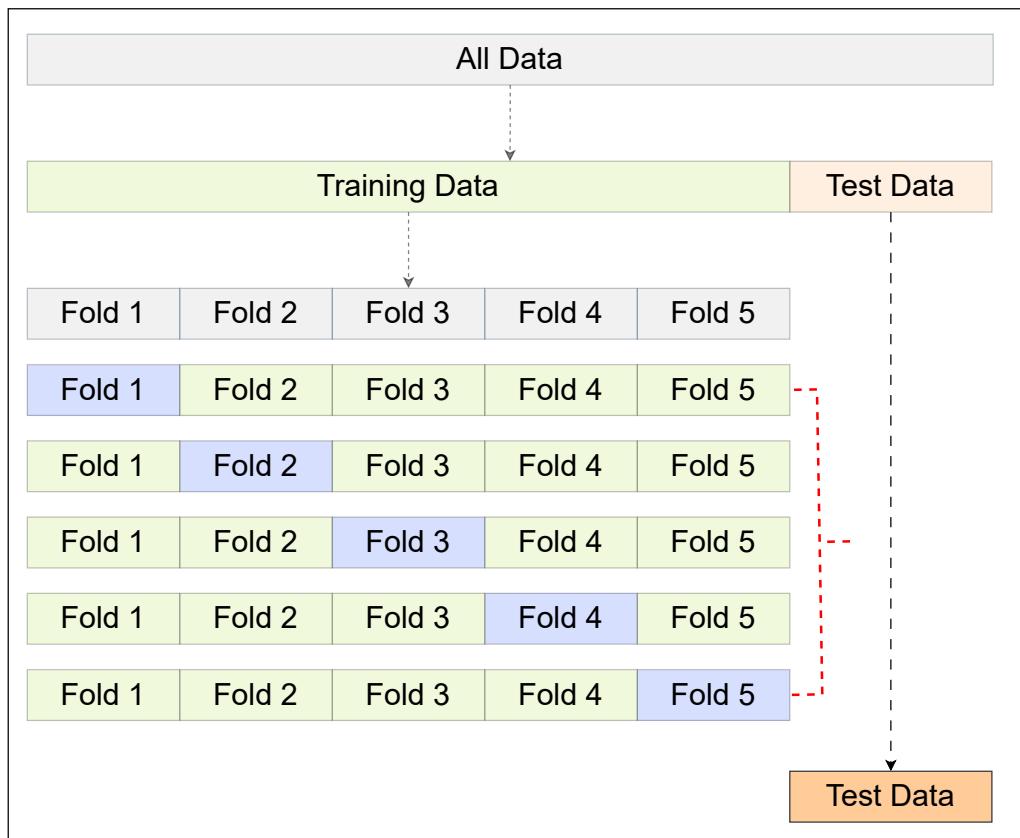
K-fold cross validation - initialise, pass $K = 5$



Perform last pass

- In the last pass, set $i = K$
- Fold $i = K$ is held back for validation
- Folds $i = 1, \dots, K - 1$ are used to train a model, giving parameters $\boldsymbol{a}^{(K)}$.
- The model is evaluated against fold $i = K$
- The distance between the actual and predicted targets is calculated and recorded as $d^{(K)}$.
- The entire training set was used to learn the “average” parameters \boldsymbol{a}
- We also have an estimate of the distribution of the prediction error d , so can compare hyperparameters

K-fold cross validation - using the results



Finding and using the best model

- Typically cross validation is used to compare effects of hyperparameters
- Use the setting(s) with the lowest cross-validation error
- Given this “best” model, we can predict the test targets
- We can then compute the prediction error on the test set, as before

K-fold cross validation

sklearn example

```
from sklearn.model_selection import cross_val_score  
# clf is some classifier, X and y are the features and target of the training set  
scores = cross_val_score(clf, X, y, cv=5)
```

- scores is a $k = 5$ element array, can be used to estimate the prediction error (or other score) while building a model
- Details of cross validation are hidden...

Featuring engineering 1: Scaling of numerical variables

Scaling - what it does

- If numeric features have different scales, e.g. [-0.005, -0.003] and [10000, 10001] some terms dominate, others are “lost”
- Better: transfer the scaling from the feature to the model parameter
- A min-max scaling is often a good choice:

$$\tilde{X} = \frac{X - X_{\min}}{X_{\max} - X_{\min}}$$

- Note that X is in the range $[X_{\min}, X_{\max}]$ but \tilde{X} is in the range $[0, 1]$.
- Other options include StandardScaler (subtract mean and divide by standard deviation) and a max-abs scaler (scales to [-1,1])

sklearn example

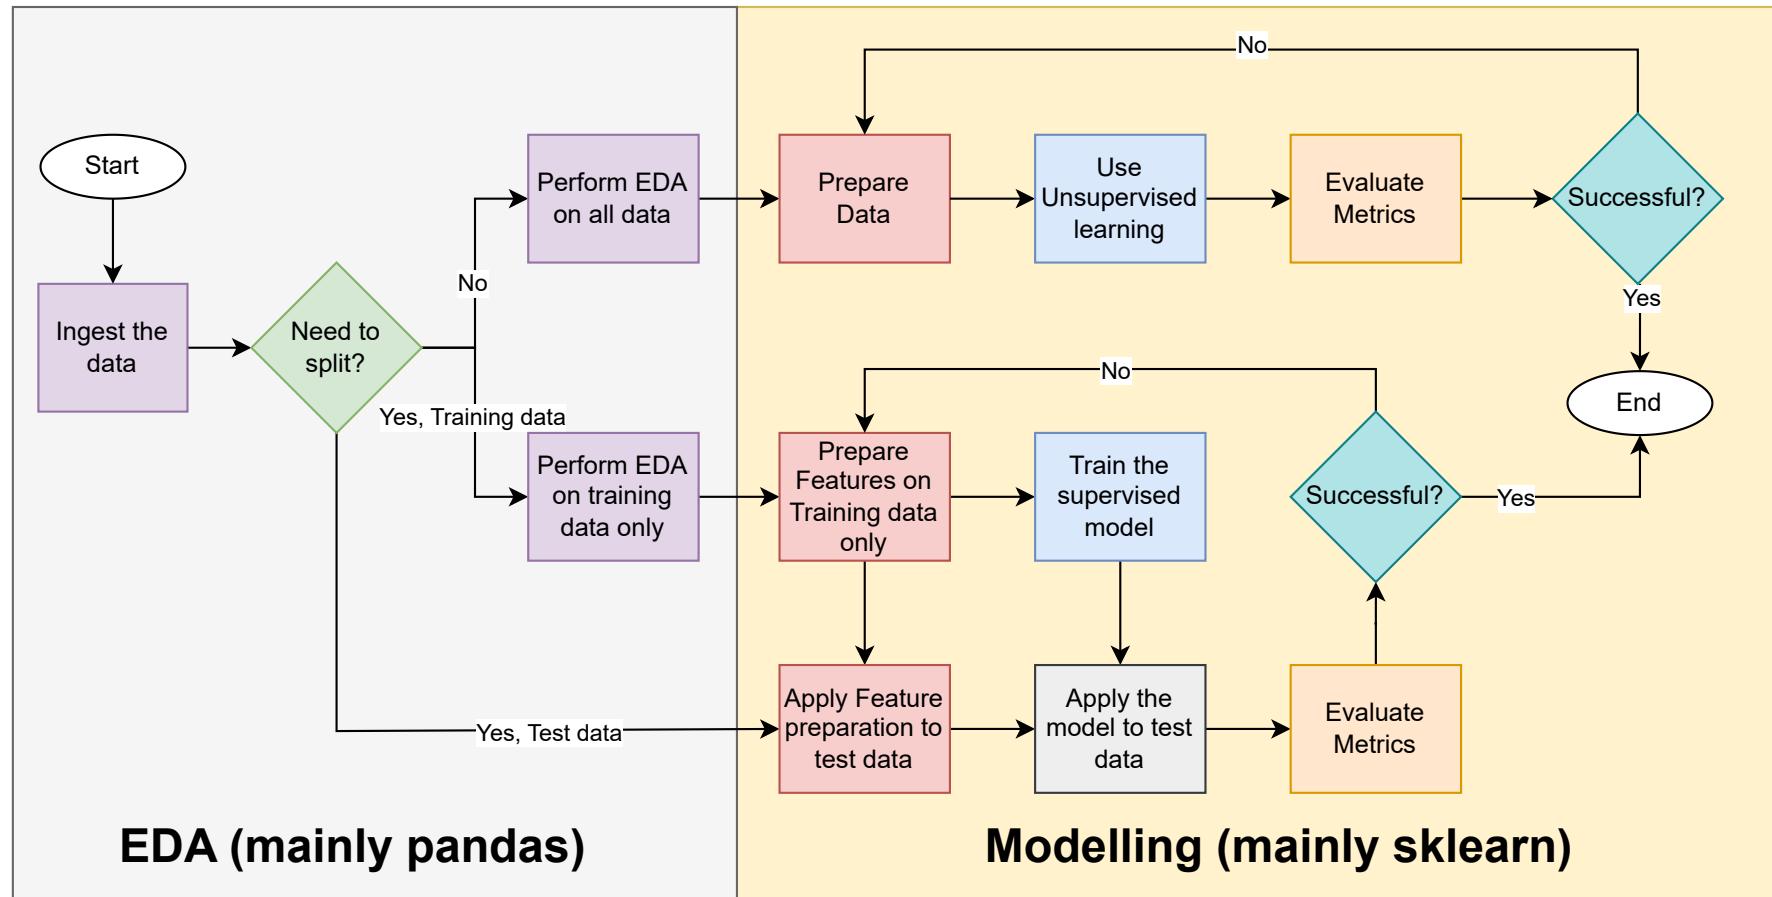
```
from sklearn.preprocessing import MinMaxScaler  
# df is a dataframe with numeric features  
scaler = MinMaxScaler()  
dfScaled = scaler.fit(df))
```

`dfScaled` can be used instead of `df` with the advantage that the fitted parameters are more accurate.

Feature Engineering 2: Choice of Features

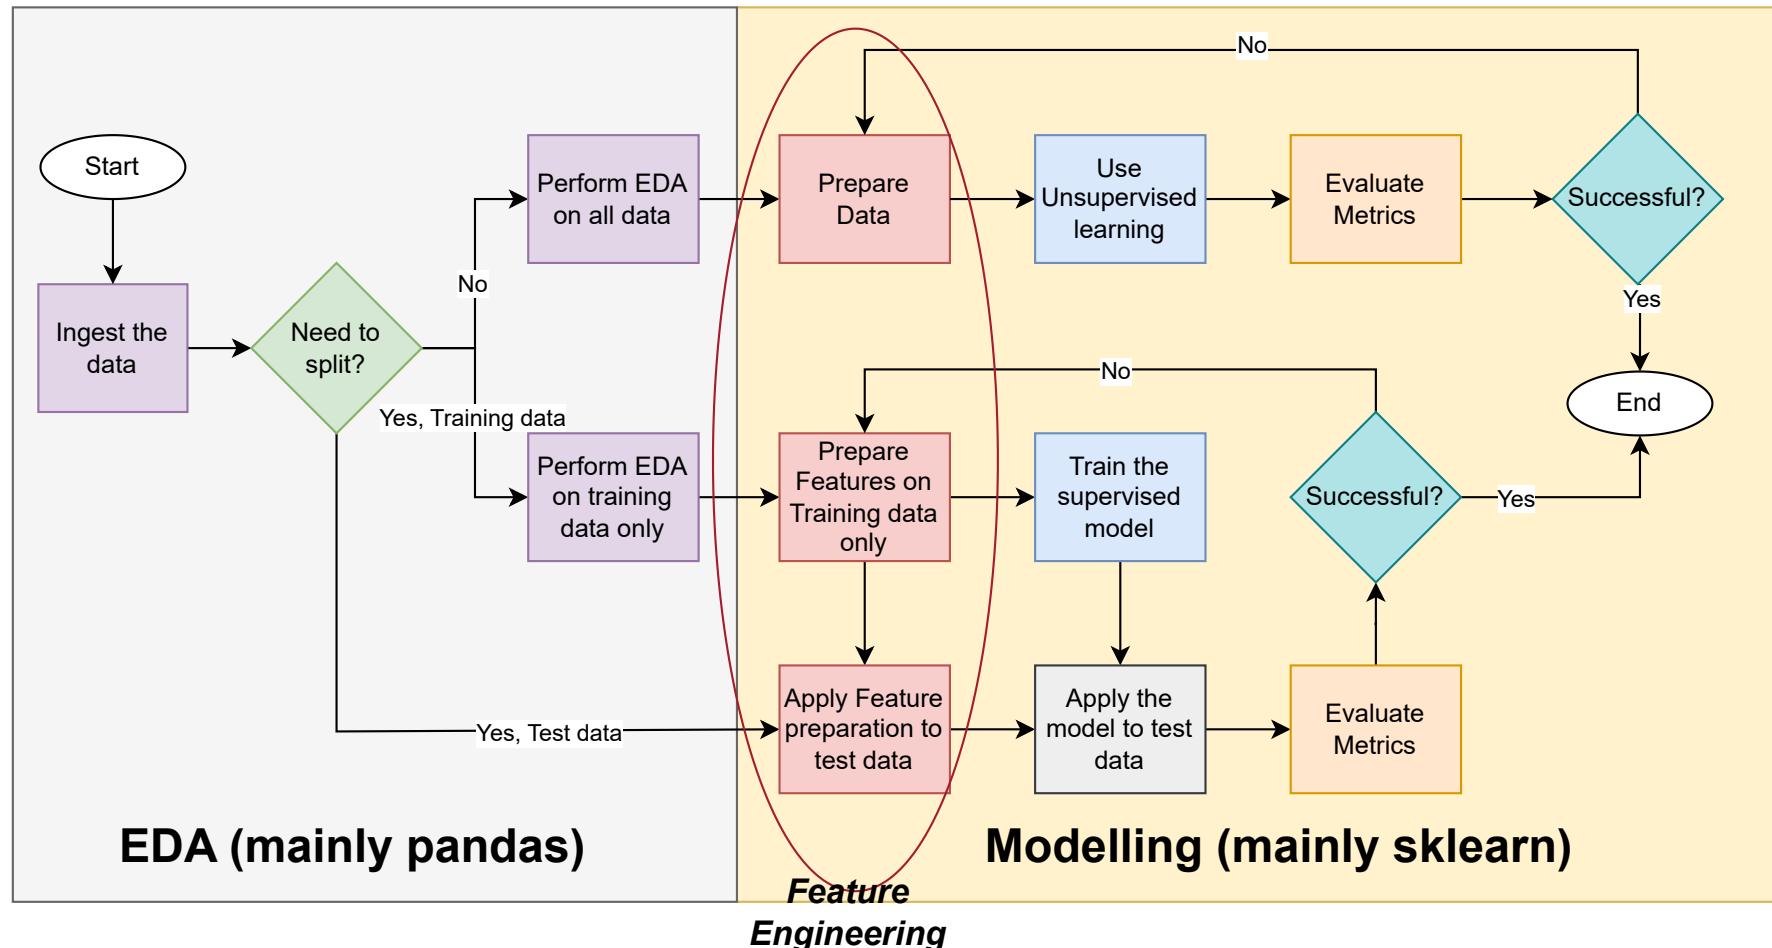
- How many to include? Use metrics to decide. Will see some when considering regression and classification.
- How do we handle different feature types? Need to encode categorical variables.
- Can we derive new numeric features? Yes, $f' = \log(f)$ etc. is possible

An overview of Machine Learning



On first glance, this might seem overwhelming, but note that the boxes are colour-coded, so that related operations share the same colour.

Where feature preparation fits...



NB: Feature preparation is informed by EDA *but is not part of EDA*. So, not part of your EDA assignment!

Feature Engineering more generally

- Use of pandas `getdummies()` requires care and is not always suitable
- Scikit-learn, imported as the `sklearn` package, supports a variety of column transformations
- Categorical columns can be features or targets, ordered or unordered
- Can also rescale numerical columns, or encode more exotic columns (other datatypes, computed columns, ...)
- As seen in the schematic, if an ML procedure is unsuccessful, more feature engineering should be considered - it can help.

Summary

- We have reviewed different types of models and considered their general form.
- We looked at the goals of modelling: minimise predictive error.
- We considered how feature engineering can help.
 - Scaling numerical features, so that variation is treated fairly between features.
 - Choosing a subset of features (more to come in future weeks...), looking for the sweet spot between under- and over-fitting.
 - Encoding categorical features as numerical dummy features (more to come in future weeks...), so they can participate in linear models
- In subsequent weeks we will put this theory into practice.

Resources

- **A Summary of the Basic Machine Learning Models**

towardsdatascience.com/a-summary-of-the-basic-machine-learning-models-e0a65627ecbe

- **Train-Test Split for Evaluating Machine Learning Algorithms**

[https://machinelearningmastery.com/
train-test-split-for-evaluating-machine-learning-algorithms](https://machinelearningmastery.com/train-test-split-for-evaluating-machine-learning-algorithms)

This week I have focused on the theory rather than its (python) implementation. This is a nice article that covers the implementation side of things.

- **Cross-Validation: Estimator Evaluator**

medium.com/swlh/cross-validation-estimator-evaluator-897d28afb4ff

Nice article that covers cross-validation in a lot more detail — we will be using many of these variants in later weeks, especially k-fold stratified.