

# Data Mining (Week 1)

dm25s1

Topic 11 : Clustering

Part 02 : Hierarchical

Preparation

Data Handling

Dr Bernard Butler

Department of Computing and Mathematics, WIT,  
[\(bernard.butler@setu.ie\)](mailto:bernard.butler@setu.ie)

Exploring Data 1

Exploring Data 2

Building Models

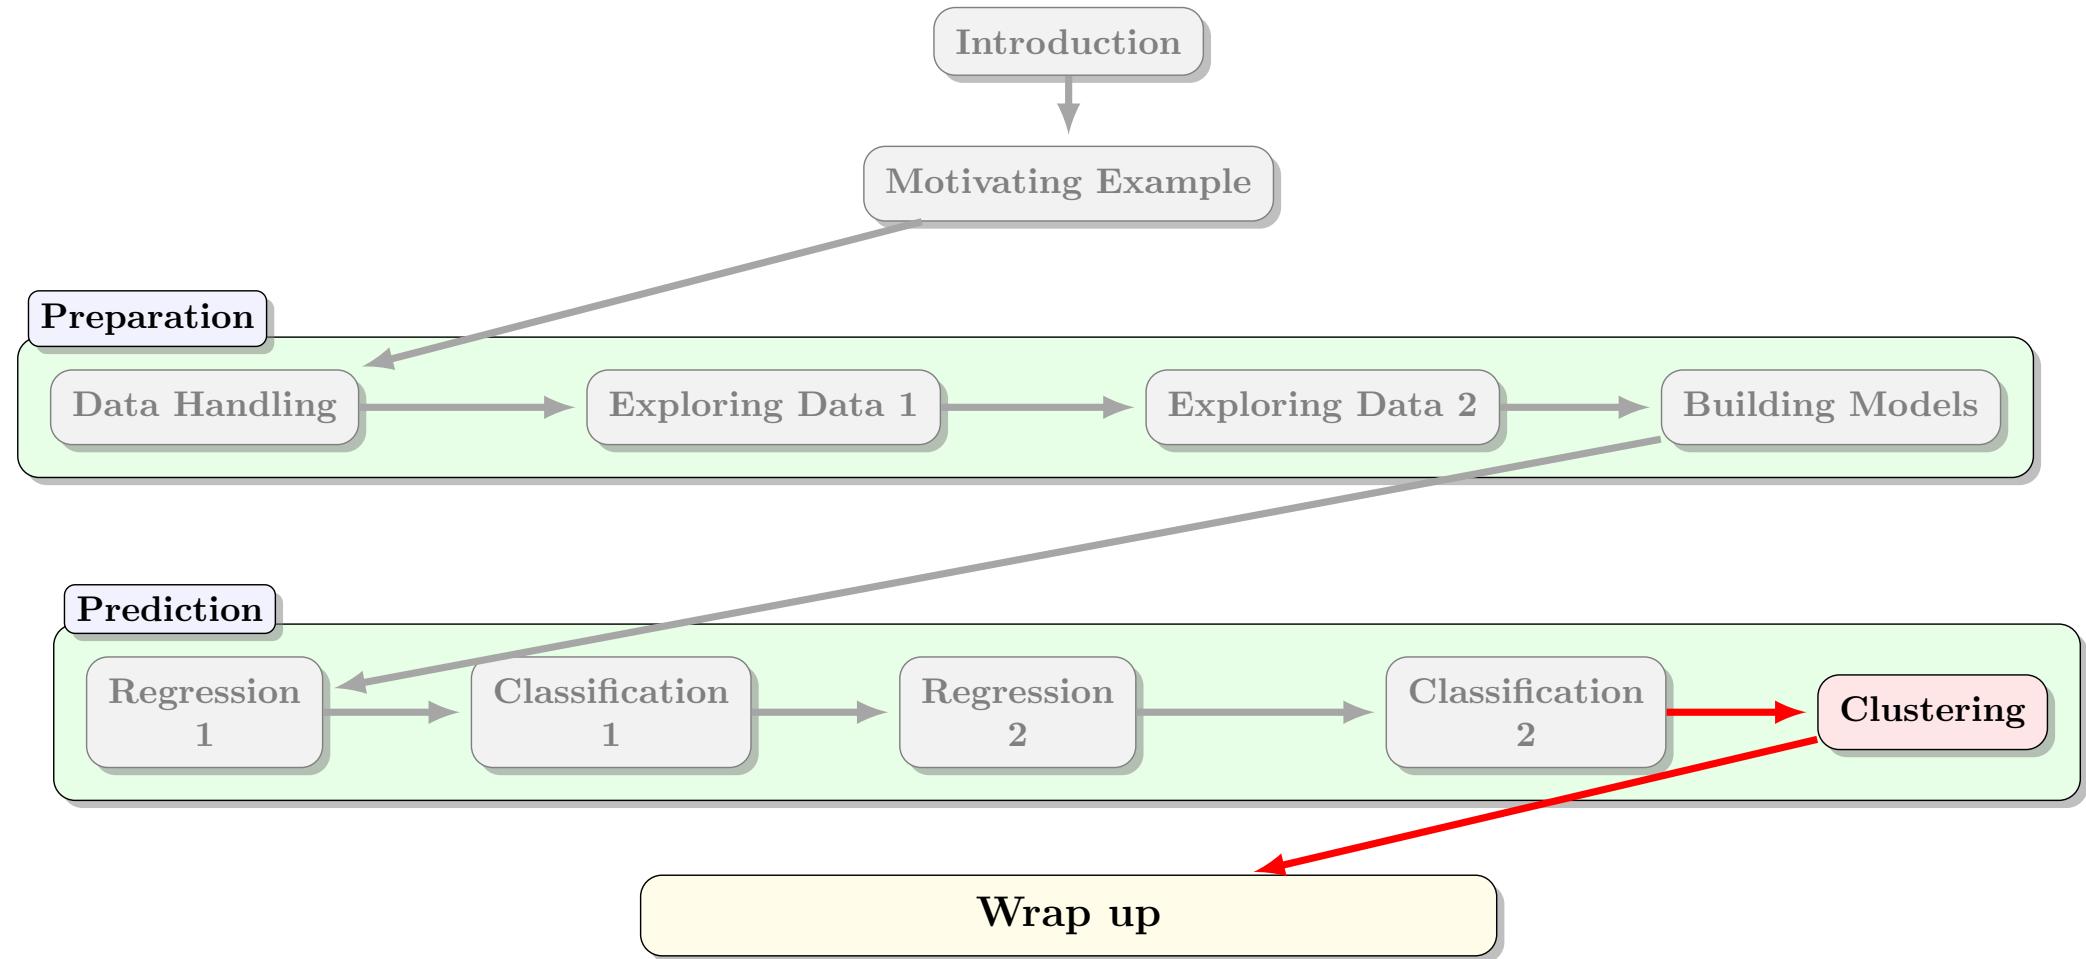
Autumn Semester, 2025

## Outline

- How to compute distances between instances
- Algorithms that partition the data

Wrap up

# Data Mining (Week 11)



# Outline

|   |    |
|---|----|
| 1. Partitioning Algorithms                    | 4  |
| 1.1. K-means                                  | 7  |
| 1.2. Soft clustering                          | 15 |
| 1.3. Expectation Maximisation (EM) iterations | 18 |
| 1.4. Density-based clustering                 | 20 |
| 1.5. Choosing $k$ for centre-based clusters   | 25 |
| 2. Review and resources                       | 31 |

# Clustering as a partitioning problem

- Often the purpose of clustering is to assign one or more labels to each observation, so that “similar” observations are given the same cluster membership label.
- In the standard case, each observation is assigned a single label, and clustering defines a (hard) *partitioning* of the data. Lloyd’s *k-means* algorithm does this.
- If each observation is assigned a membership probability for each cluster, this is a *soft partitioning* of the data. A hard partition can be derived by choosing, for each observation, the cluster for which it has the highest probability of membership. *Gaussian Mixture* models can be used for this purpose.
- Some clustering algorithms, notably *density-based clustering*, do not always assign a label to each observation. However, if an observation is assigned a label, it will be just one such label.

## Definition 1 (Representation-based clustering)

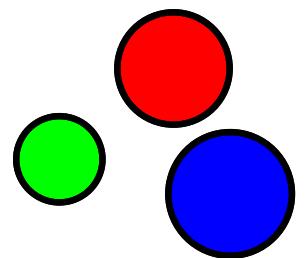
... finds a region around a *cluster centre* so that observations can be assigned to the cluster if they are found in that region.

## Definition 2 (Density-based clustering)

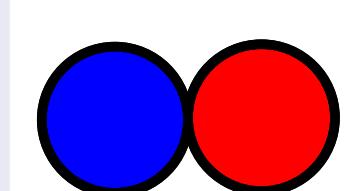
... looks for regions, possibly non-convex, where the data density is higher, and assigns observations in those regions to the relevant cluster. Any other observations are assumed to be either “noise” or “border” observations.

# Types of partitional clustering

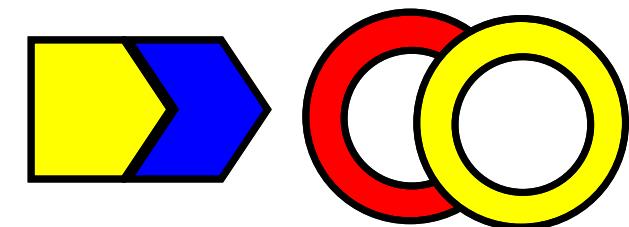
Well-separated clusters



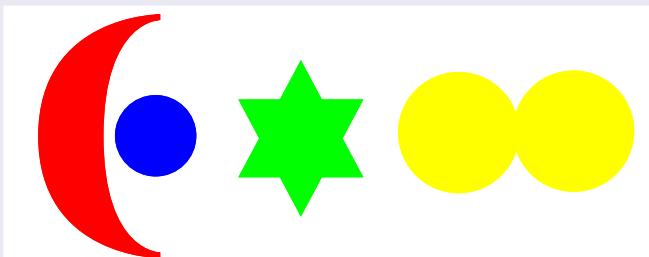
Centre-based clusters



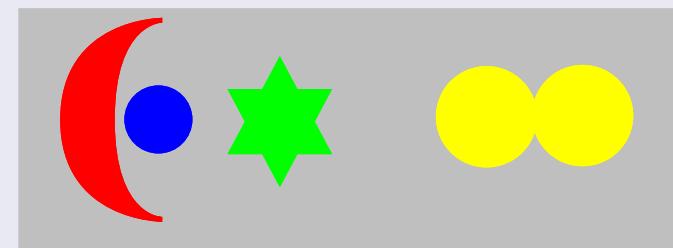
Conceptual clusters



Contiguity-based clusters



Density-based clusters



# K-means algorithm: Overview

- The  $k$ -means algorithm assigns each observation to one of  $k$  clusters, by finding the nearest cluster centre for that observation (E-step).
- Each cluster centre is calculated as the centroid of the observations assigned to that cluster (M-step).
- The algorithms proceeds in two steps (E-M). At each iteration, the algorithm finds the nearest centre for each observation, then assigns it to that centre and recomputes the centres.
- Lloyd's algorithm is an example EM-algorithm: Expectation-Maximisation (general algorithm, used in many scenarios, especially clustering).
- Variants include

**Mini-batch k-means** : work with a random sample of the data at each iteration: scales better, small loss of accuracy

**kmeans++** : Choose initial centres that are well-separated from each other; “normal” k-means afterwards.

**k-medoids** : Manhattan ( $\ell_1$ ) distance is used instead of Euclidean ( $\ell_2$ ), and centres are constrained to be data points); PAM and CLARA algorithms.

- Generally Lloyd's algorithm is robust, although it is affected by the choice of initial centres, and care must be taken to avoid empty clusters

# K-means algorithm: Detail

## Method (k-means algorithm)

$t \leftarrow 0;$

Initialise centres  $\{\mu_j^t, j = 1, \dots, k\}$ : choose  $k$  points randomly, without replacement;

**repeat**

$t \leftarrow t + 1;$

$C_j \leftarrow \emptyset, \forall j = 1, \dots, k;$

**for all**  $x$  **do**

▷ Assign  $x_j$  to the nearest centroid from the previous iteration

$j^* \leftarrow \arg \min_i \left\{ \|x_j - \mu_i^{t-1}\|^2 \right\};$

$C_{j^*} \leftarrow C_{j^*} \cup \{x_j\};$

**end for**

▷ Centroid Update Step M

**for all**  $i = 1$  to  $k$  **do**

$\mu_i^t \leftarrow \frac{1}{|C_i|} \sum_{x_j \in C_i} x_j;$

**end for**

**until**  $\sum_{i=1}^k \|\mu_i^t - \mu_i^{t-1}\|^2 \leq \epsilon$

The termination condition is that the difference in centre positions should not exceed a small tolerance  $\epsilon$ . This happens when points stay in their cluster from iteration  $p$  to  $p + 1$ , so cluster centre stays same.

# Clustering categorical data

- k-means uses **centroids** and **Euclidean distance**
- So it cannot be applied directly to categorical data...
- Options
  - **Either** encode categorical columns as integers - can now compute distances
  - Since this can dramatically increase the dimensionality, dimensionality reduction (e.g., PCA) might be needed
  - **Or** Use k-modes on the original data if *all* the data is categorical
  - **Or** Use k-prototypes on the original data if some data is categorical and some is numerical

# “K-means” for categorical data: k-modes

k-means uses centroids and Euclidean distance; k-modes uses modes and Hamming distance

Strengths and weaknesses: mostly similar to k-means

- Guaranteed to converge (eventually): helped by good choice of  $k$  and initial modes
- Iterates to a local minimum: result quality depends on initial modes
- Distances are integers: need to choose between tied distances when assigning cases to clusters

## Implementation - Setup

Installation:

```
conda install conda-forge::kmodes
```

Imports:

```
from kmodes.kmodes import KModes
import pandas as pd
import numpy as np
```

Since k-modes, like most EM algorithms, starts from a random cluster assignment and iterates to improve it, a poor initial choice can result in sub-optimal results. Here, we ask it to start from 4 cluster assignment choices, to iterate to completion for each, and to pick the best overall cluster assignment.

```
model=KModes(n_clusters=3, random_state=42, n_init=4)
```

# k-modes in practice

## Deriving cluster centroids

```
fittedModel=model.fit(df)  
print("Cluster centroids – archetypal student grades")  
print(fittedModel.cluster_centroids_)
```

Cluster centroids – archetypal student grades

```
[['A' 'A' 'B' 'B' 'A']  
 ['A' 'C' 'B' 'A' 'C']  
 ['A' 'B' 'A' 'C' 'B']]
```

- Note these centroids indicate a “typical student” in that cluster, but this does not need to match any of the students that were used when learning the cluster assignment.
- That is, the individual grades are the modes of the grades in that cluster, but that does not mean that a student in the cluster has that combination of grades.

## k-modes assigns each student to a (0,1,2) cluster

```
clusters = fittedModel.predict(df)
df["ClusterID"] = clusters
print("Allocation of students to clusters")
print(df)
```

### Allocation of students to clusters

| Student           | English | Maths | History | Geography | Science | ClusterID |
|-------------------|---------|-------|---------|-----------|---------|-----------|
| Beryl Smart       | A       | B     | A       | B         | A       | 0         |
| Sydney Whitworth  | C       | C     | B       | A         | C       | 1         |
| Nora Waite        | C       | A     | B       | B         | A       | 0         |
| Carrie Aldridge   | B       | A     | A       | B         | C       | 0         |
| Ravinder Townsend | A       | B     | B       | A         | C       | 1         |
| Antonio Hunter    | B       | A     | C       | C         | C       | 0         |
| Clive Sheldon     | A       | A     | A       | A         | A       | 0         |
| Lynette England   | A       | C     | B       | B         | B       | 0         |
| Hilary Farrow     | A       | B     | B       | A         | A       | 0         |
| Cristina Rogers   | C       | C     | D       | B         | A       | 0         |
| Nana Gilbert      | A       | C     | B       | B         | C       | 1         |
| Miriam Moore      | A       | B     | A       | C         | B       | 2         |
| Tara Lomas        | B       | C     | C       | D         | B       | 1         |
| Sam Humphrey      | A       | B     | B       | B         | B       | 0         |
| Olga Pickles      | A       | A     | B       | B         | A       | 0         |

# “K-means” for numerical *and* categorical data: k-prototypes

Combine k-means (on numerical data) and k-modes (on categorical data) in one clustering algorithm.

## Intuition

- A **prototype** instance has representative values of numerical and categorical features.
- Distance is a linear combination of the Euclidean (numerical) and Hamming (categorical) distances
- k-prototypes is similar to k-means or k-modes, with similar strengths and weaknesses

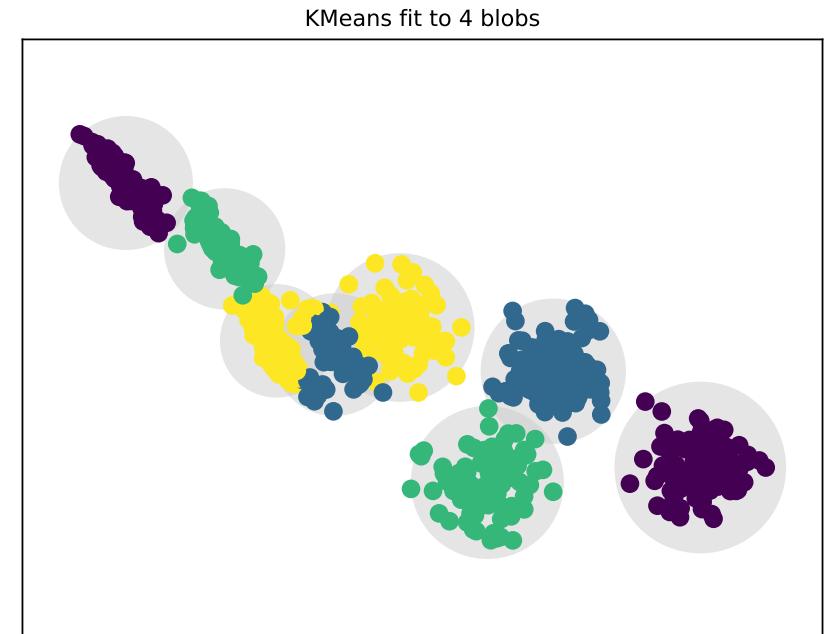
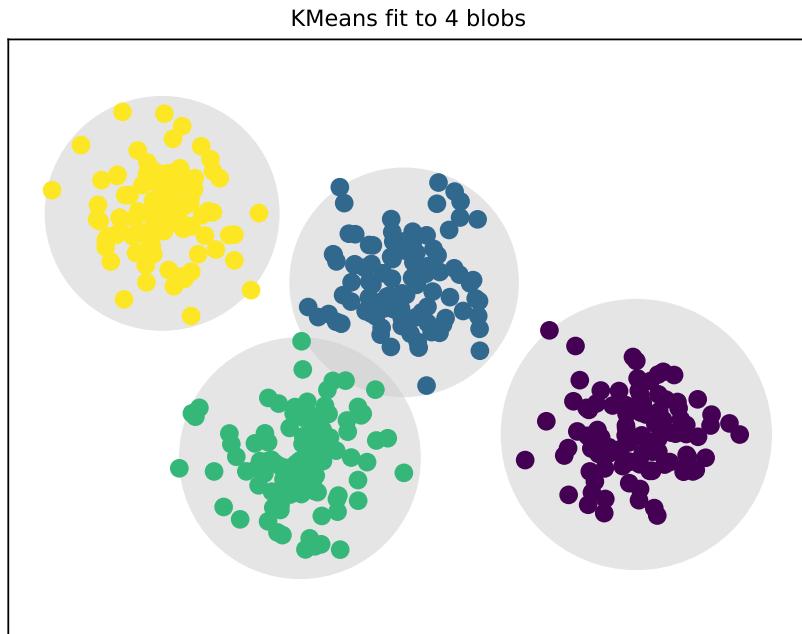
## Implementation

```

kRange = range(1,8)
allCols = numCols + allCatCols
catColIDs = list(range(len(numCols),len(numCols)+len(allCatCols)))
scores = dict()
for k in kRange:
    # Use Huang initialisation, use 5 random starting starting points, turn off logging
    model = KPrototypes(n_clusters=k, init='Huang', verbose=0, random_state=42, n_init=5)
    # Note that we need to tell the model which are the categorical columns
    fittedModel = model.fit(df[allCols], categorical=catColIDs)
    scores[k] = fittedModel.cost_
print(scores)

```

## K-means algorithm: In practice



With the original globular clusters,  $k$ -means was able to find the centres and clusters easily.

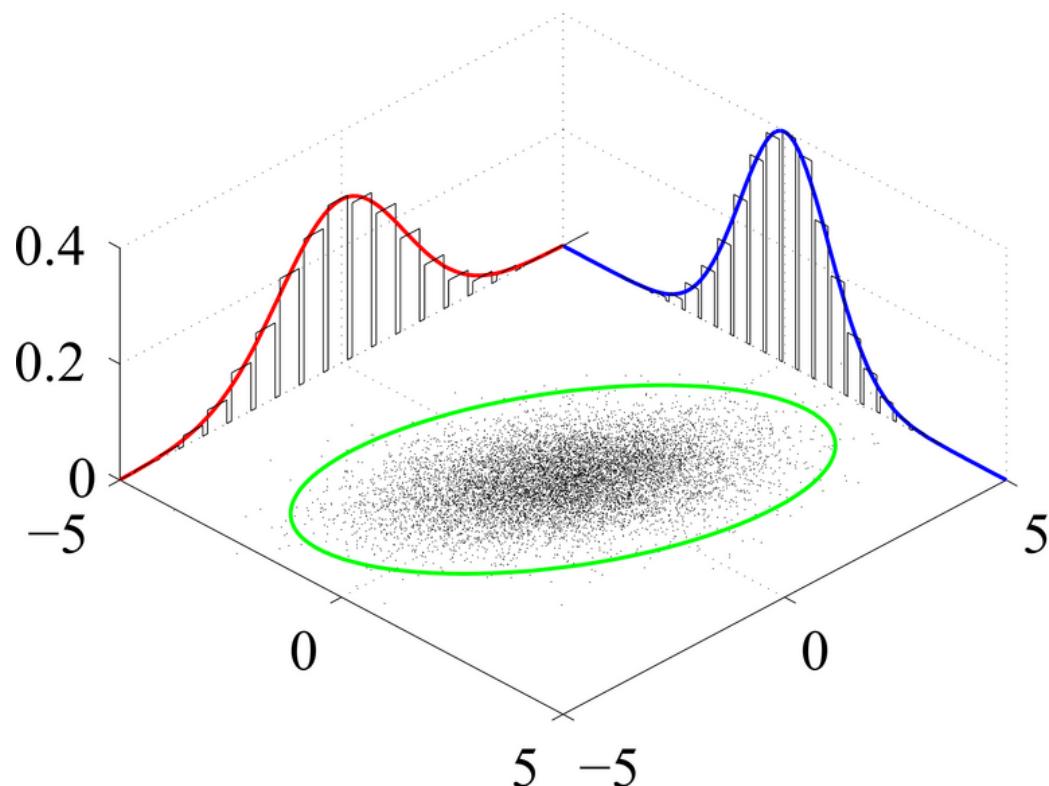
$k$ -means minimises the within-cluster sum of squared distances (also known as *inertia*) so the choice of distance function is critical.

With the stretched clusters,  $k$ -means had more difficulty, e.g., with the yellow and purple clusters.

# Probabilistic models for clustering

- k-means is an example of *hard* clustering, where each data point is mapped to a single cluster
- As such it is best suited to well-separated clusters - but what if they are close or even overlap?
- *fuzzy* clustering: points are assigned to multiple clusters, and are given a membership score in [0,1] for each cluster
- fuzzy c-means algorithm is a straightforward extension of k-means, just using probability  $P(x_i, \mu_j)$  to weight each point  $x_i$  when calculating the centroid of each cluster  $\mu_j$  (M-step)
- $P$  is a function of the relative Euclidean distances to the cluster centres  $\{\mu_j\}$
- This probability function can be generalised, notably to take account of the *shape* of the clusters and not just their centres, leading to *Gaussian Mixture Model* (GMM) probabilistic clustering

## Review: Multivariate (2D) Gaussian/Normal distribution



- The distribution can have different dimensions that do not need to align with the coordinate axes; captured as a  $2 \times 2$  covariance matrix  $C$
- The distribution stretches to infinity in the plane, but points far from the centre of the distribution have very low probability.
- A collection of clusters can be modelled by overlaying a *mixture* of such Gaussian distributions on the plane.

Source: Wikipedia

# Review of Bayes Theorem

## Use of Bayes Theorem in Classification

**Likelihood** is the probability of the data given the label. **Prior** measures our belief about how likely each label is *before* we have seen any data. The **Posterior** includes influences of both the Prior and the Likelihood.

$$P(y = c|x) = \frac{P(x|y = c)P(y = c)}{P(x)}$$

The Posterior here is  $P(y = c|x)$ , the Likelihood is  $P(x|y = c)$  and the Prior is  $P(y = c)$ .  $P(x)$  is a normalizing constant that measures how likely the observed data  $x$  is.

When used for Gaussian Mixture Models, there is not just a *single* cluster label  $c$ , but a linear combination of many.

## Overview of EM algorithm for GMM clustering

**E-step** For each  $x_i$ , calculate the probability that  $x_i$  belongs to the  $j^{\text{th}}$  distribution

$$P(\Theta_j|x_i, \Theta) = \frac{P(x_i|\Theta_j)}{\sum_{l=1}^k P(x_i|\Theta_l)},$$

where  $\Theta_j$  is the set of parameters defining Gaussian distribution  $j$ , namely its centre  $\mu_j$  and covariance matrix  $C_j$ .

**M-step** Maximise the expected likelihood  $P(\{x_i\}|\Theta)$  by updating the Gaussian mixture. That is, for each  $\mu_j$  and  $C_j$ , use all  $x_i$  and the  $P(\Theta_j|x_i, \Theta)$  computed in the E-step to derive the new Gaussian distribution parameters.

Note that the E-step computes a membership probability for each point based on all the Gaussian models and their parameters.

By contrast, the M-step computes the new Gaussian models based on all the points and their membership probabilities.

Lloyd's k-means algorithm is equivalent: the membership probability is either 1 (allocated to this cluster) or 0 (not allocated to this cluster) for each point. The M-step re-computes the cluster centres based on all the points and their cluster assignment.

# GMM compared with k-means

|         | k-means  | GMM  |
|---------|--|--|
| E-step  | Compute membership probability which is either 1 (allocated to this cluster) or 0 (not allocated to this cluster) for each point | Compute membership probability for each point based on all the Gaussian models and their parameters. |
| M-step  | Recompute the new cluster centres based on all the points and their cluster assignment   | Recompute the new Gaussian models based on all the points and their membership probabilities         |
| Use for | Well-separated   | Centre-based or well-separated   |
| Shape   | nondirectional (“spherical”)   | directional (“ellipsoidal”) or nondirectional  |

## Relaxing the constraints: density-based clustering

k-means and GMM are both characterised by the following properties:

- the number of clusters  $k$  must be specified beforehand
- clusters have a convex shape
- they work best when the clusters are linearly separable
- all points are assigned to clusters, so can be sensitive to outliers

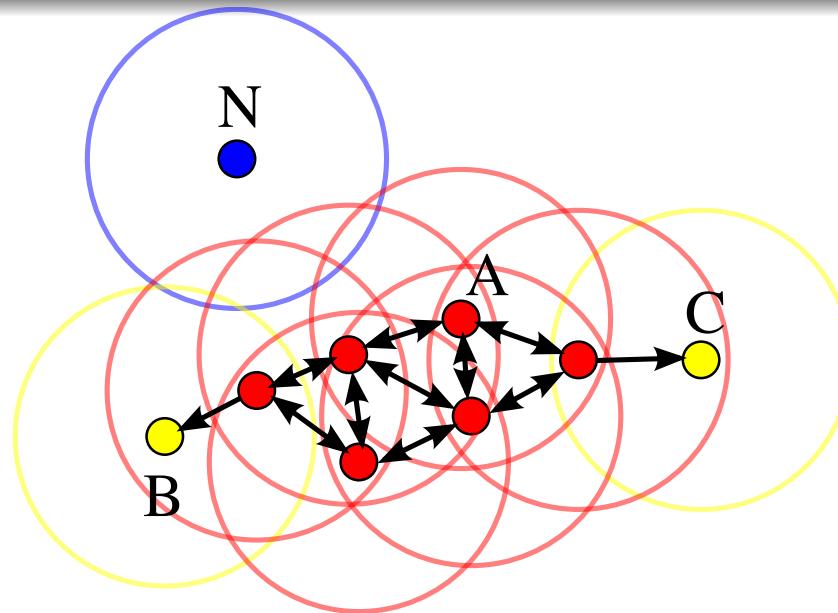
➤ Density-based clustering relaxes these conditions

It uses the heuristic that clusters are (arbitrarily-shaped) contiguous regions with high datapoint density.

Datapoints outside these regions represent *noise* and are ignored.

Rather than specifying  $k$ , the user specifies density thresholds.

# Relaxing the constraints: density-based clustering



Source: wikipedia

- A, B and C are directly connected points.
- A is a **core** point
- B and C are **border** points.
- N is a **noise** point and so is not assigned to a cluster.
- The connected component of the 8 points (6 red, 2 yellow; including A,B,C) forms a cluster.

# DBSCAN algorithm and its concepts

## Definition 3 (DBSCAN)

Density-Based Spatial Clustering of Applications with Noise (DBSCAN): an algorithm for deriving clusters in areas of high data density.

## Definition 4 (eps-neighbourhood)

Epsilon  $\epsilon$  parameter defines a region of points  $t$  around a point  $x$  where  $\|t - x\| < \epsilon$ .

## Definition 5 (core point)

Point with at least  $\text{MinPts}-1$  other points in its eps-neighbourhood.

## Definition 6 (border point)

Point with less than  $\text{MinPts}-1$  other points in its eps-neighbourhood, but at least one is a core point.

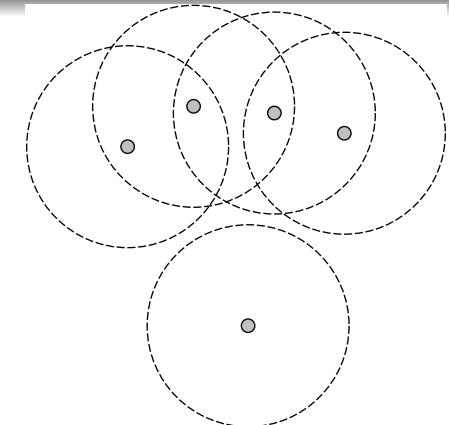
## Definition 7 (noise point)

Point with less than  $\text{MinPts}-1$  other non-core points in its eps-neighbourhood.

# Development of the algorithm

## Definition 8 (Direct density reachable)

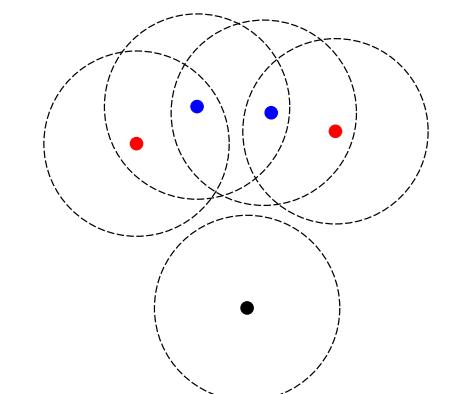
Point  $x_A$  is directly density reachable from  $x_B$  iff  $x_A$  is in the  $\text{eps}$ -neighborhood of  $x_B$  and  $x_B$  is a *core point*.



Core, border and noise points are coloured blue, red and black below.

## Definition 9 (Density reachable)

Point  $x_A$  is density reachable from  $x_B$  if there is a set of core points in each other's  $\text{eps}$ -neighbourhood between  $x_A$  and  $x_B$ .



## Definition 10 (Density connected)

Points  $x_A$  and  $x_B$  are density connected if there exists a core point  $x_C$  so that both  $x_A$  and  $x_B$  are density reachable from  $x_C$ .

# Steps of the DBSCAN algorithm

## Method (DBSCAN)

- 1 Find the  $\epsilon$  neighbors of every point.
  - 2 Identify the core points with more than minPts neighbors.
  - 3 Derive the *connected component* graphs of core points, assigning edges between core points that are less than  $\epsilon$  apart.
  - 4 Identify the border points and assign them to their nearest cluster.
  - 5 Label any remaining points as *noise*.
- A variant (HDBSCAN) excludes border points from the cluster, treating them as noise points (can be more robust).
  - Another variant (OPTICS) places the points in a priority queue, ordered by reachability distance (updating is slower, but handles varying density better).

# Choosing $k$ , the number of clusters

## ➤ How can we decide on $k$ for k-means and GMM?

- We can do this *graphically* (plot clusters for each  $k$ ) or by using *scores*.
- Plot within-cluster sum of squared distances (inertia) against  $k$  and look for  $k$  at the “elbow”.
- Use `kmeans.inertia_` as the score for a given instance of the `kmeans` classifier.
- Can also compute inertia for other partitional clustering techniques, but this is more work and interpretation is more difficult.

## Silhouette scores - derivation

➤ How much is any point in a cluster nearer its peers than it is to points in the nearest of the other clusters?

### Method (Silhouette score)

**Require:** Clustering where the  $i$  point is assigned to cluster  $C(i)$  and there are  $k$  such clusters

**for all** point  $i$  in cluster  $C(i)$  **do**

    Calculate  $a(i)$ , the mean distance between  $i$  and all the other points in  $C(i)$ .  $\triangleright a(i) \equiv 0$  is there is no other point in  $C(i)$ .

    Calculate  $b(i)$ , minimum of the mean distances between  $i$  and all the other points in each of  $C(j)$  where  $j \neq i$ .

    Silhouette  $s(i) = 1 - a(i)/b(i)$  if  $a(i) < b(i)$ ,  $s(i) = 0$  if  $a(i) = b(i)$  and  $s(i) = b(i)/a(i) - 1$  if  $a(i) > b(i)$ .

**end for**

The mean of  $s(i)$  over all points ( $\bar{s}_k$ ) is a measure of the clustering efficiency for that value of  $k$ .

The  $k$  associated with the *maximum* of these  $\bar{s}_k$  silhouette scores is the best choice of  $k$ .

There are many other scores but they require more advanced mathematics and are out of scope for this module.

# Silhouette scores - examples

*Code to compute the silhouette score*

```
from sklearn.metrics import silhouette_samples, silhouette_score

k = 4
clusterer = KMeans(n_clusters=k, random_state=10)
cluster_labels = clusterer.fit_predict(X)

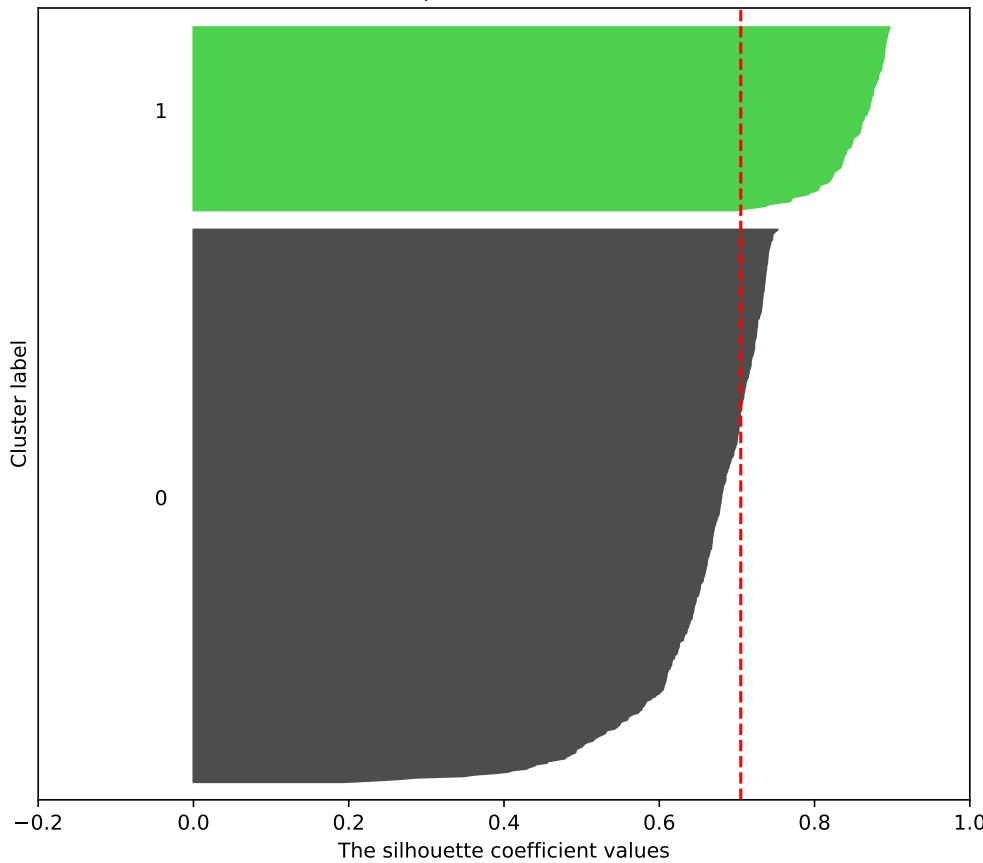
# silhouette_score is the average silhouette for all the samples
silhouette_avg = silhouette_score(X, cluster_labels)
print(silhouette_avg)
```

0.6505186632729437

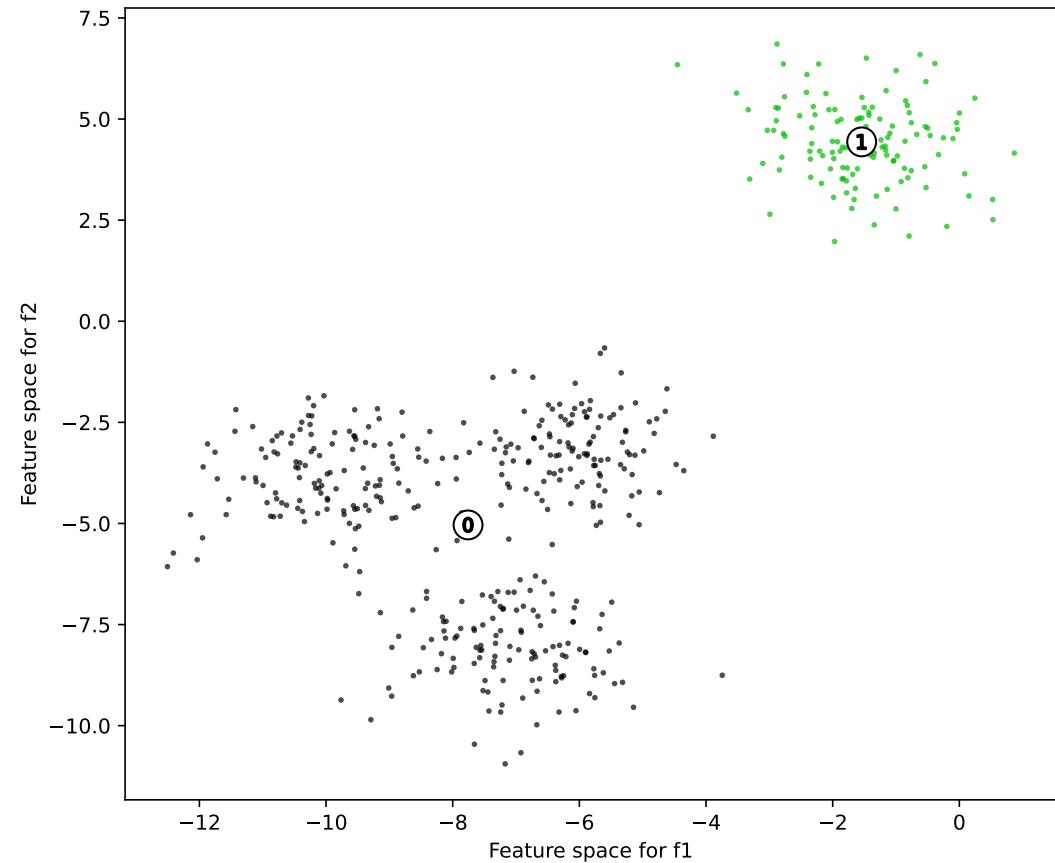
# Silhouette score with $k = 2$ - looking good

**Silhouette analysis for KMeans clustering on sample data with n\_clusters = 2**

The silhouette plot for the various clusters.



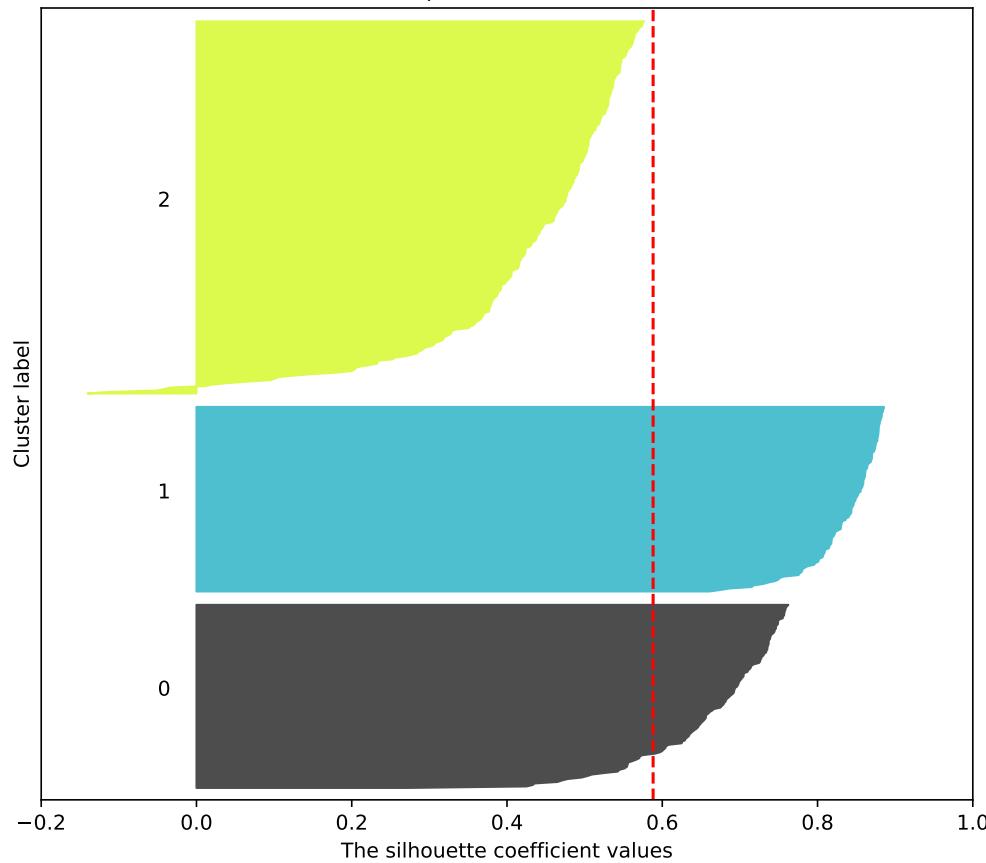
The visualization of the clustered data.



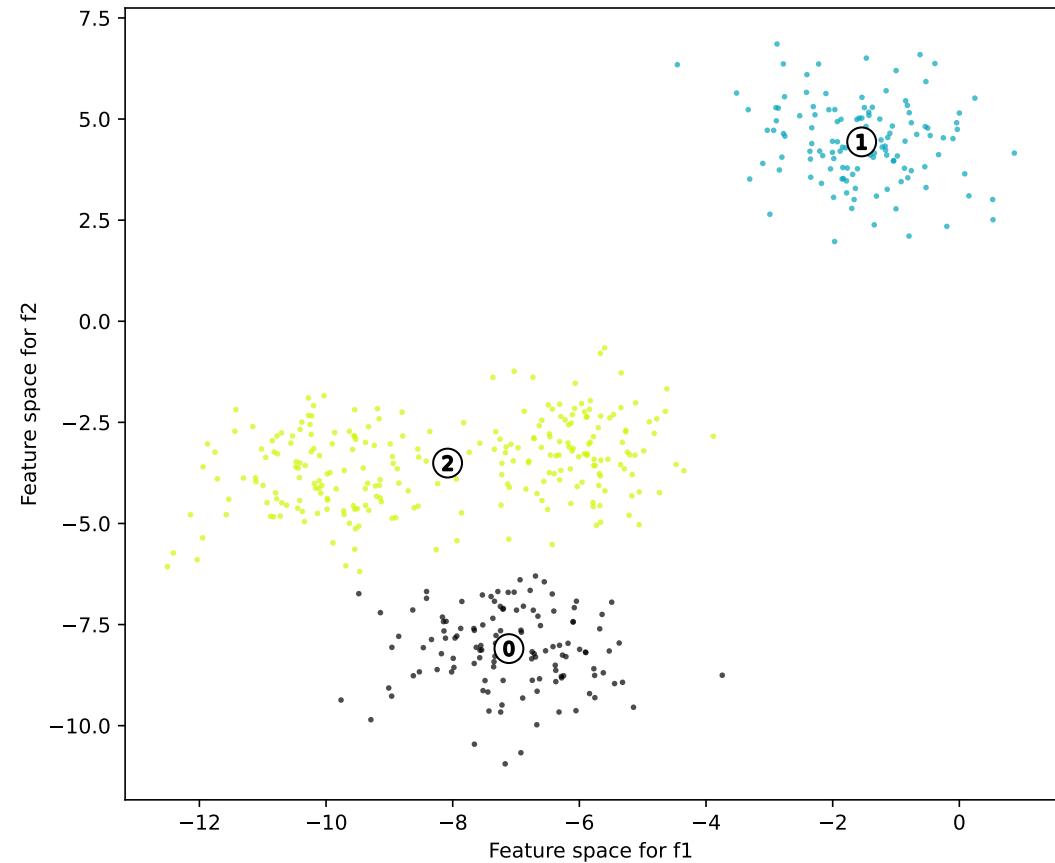
# Silhouette score with $k = 3$ - not looking good

**Silhouette analysis for KMeans clustering on sample data with n\_clusters = 3**

The silhouette plot for the various clusters.



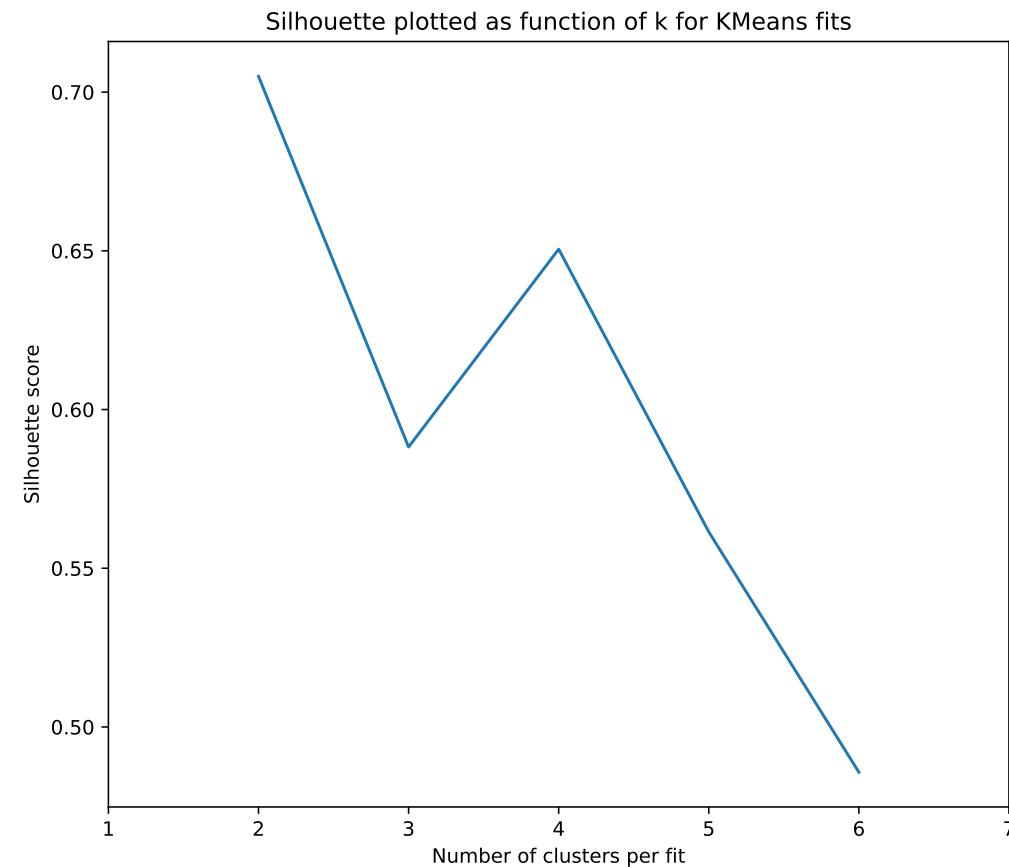
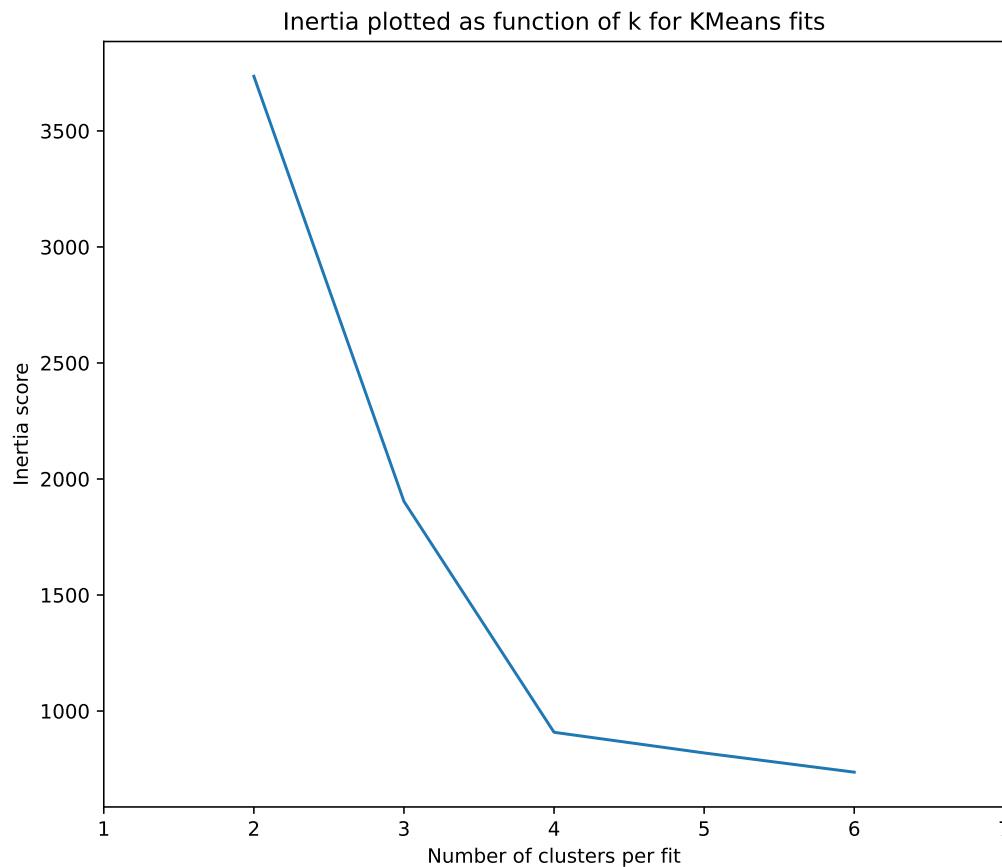
The visualization of the clustered data.



# Comparing both scoring systems

*Inertia/elbow plot and silhouette plot on the same data*

**Comparison of inertia and silhouette scores for estimating k**



# Choice of hyperparameters

**AGNES** : choose distance function and linkage. Usually Ward or single linkage work best.

**k-means, etc.** : choose distance function, starting condition (cf kmeans++), aggregation (cf k-medoids),  $k$

**GMM** : choose distance function,  $k$

**DBSCAN** : choose distance function, minpts, eps

## Tips

- Good idea to scale so that clusters are approximately (hyper)spherical.
- Can be good idea to transform data before clustering.
- Dendrogram is good for visualising structure when data is more than 3-D.

# Summary

- Clustering is perhaps the best known form of unsupervised learning
- Hierarchical clustering can provide insights into the structure of a data set - very useful when exploring data for other techniques
- Partitional clustering can be used to label points according to which cluster they belong to
- Partitional classification has many approaches: centre-based and density based are most common
- Clustering can be used to help create training data for classification purposes (c.f., the digits notebook used in the practical)